

# Recidivism\_Predict

February 11, 2023

```
[ ]: #Christine Orosco, Myra Wheeler, Kyle Hansen
#Project to predict probability of recidivism for a group of inmates
# Dataset is the 3-yr Recidivism Rates for offenders released from Iowa prisons
```

```
[1]: # Import libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[2]: # Import dataset
df = pd.read_csv("3-Year_Recidivism_for_Offenders_Released_from_Prison_in_Iowa.
↪csv")
df.head()
```

```
[2]:
```

	Fiscal Year Released	Recidivism Reporting Year	Main Supervising District	\
0	2010	2013	7JD	
1	2010	2013	NaN	
2	2010	2013	5JD	
3	2010	2013	6JD	
4	2010	2013	NaN	

	Release Type	Race - Ethnicity	Age At Release	Sex	\
0	Parole	Black - Non-Hispanic	25-34	Male	
1	Discharged - End of Sentence	White - Non-Hispanic	25-34	Male	
2	Parole	White - Non-Hispanic	35-44	Male	
3	Parole	White - Non-Hispanic	25-34	Male	
4	Discharged - End of Sentence	Black - Non-Hispanic	35-44	Male	

	Offense Classification	Offense Type	Offense Subtype	Return to Prison	\
0	C Felony	Violent	Robbery	Yes	
1	D Felony	Property	Theft	Yes	
2	B Felony	Drug	Trafficking	Yes	
3	B Felony	Other	Other Criminal	No	
4	D Felony	Violent	Assault	Yes	

	Days to Return	Recidivism Type	New Offense Classification	New Offense Type	\
--	----------------	-----------------	----------------------------	------------------	---

0	433.0	New	C Felony	Drug
1	453.0	Tech	NaN	NaN
2	832.0	Tech	NaN	NaN
3	NaN	No Recidivism	NaN	NaN
4	116.0	Tech	NaN	NaN

	New Offense Sub Type	Target Population
0	Trafficking	Yes
1	NaN	No
2	NaN	Yes
3	NaN	Yes
4	NaN	No

```
[3]: # Get a list of column names
list(df.columns)
```

```
[3]: ['Fiscal Year Released',
      'Recidivism Reporting Year',
      'Main Supervising District',
      'Release Type',
      'Race - Ethnicity',
      'Age At Release ',
      'Sex',
      'Offense Classification',
      'Offense Type',
      'Offense Subtype',
      'Return to Prison',
      'Days to Return',
      'Recidivism Type',
      'New Offense Classification',
      'New Offense Type',
      'New Offense Sub Type',
      'Target Population']
```

```
[3]: # Pull out the info for the target variable to plot
recid_counts = pd.DataFrame(df['Return to Prison'].value_counts())
recid_counts.reset_index(inplace=True)
recid_counts
```

```
[3]:   index  Return to Prison
0     No             17339
1    Yes              8681
```

```
[78]: # Plot counts of target variable
fig, ax = plt.subplots()

# Remove the right border
```

```

right = ax.spines["right"]
right.set_visible(False)

# Remove the top border
top = ax.spines["top"]
top.set_visible(False)

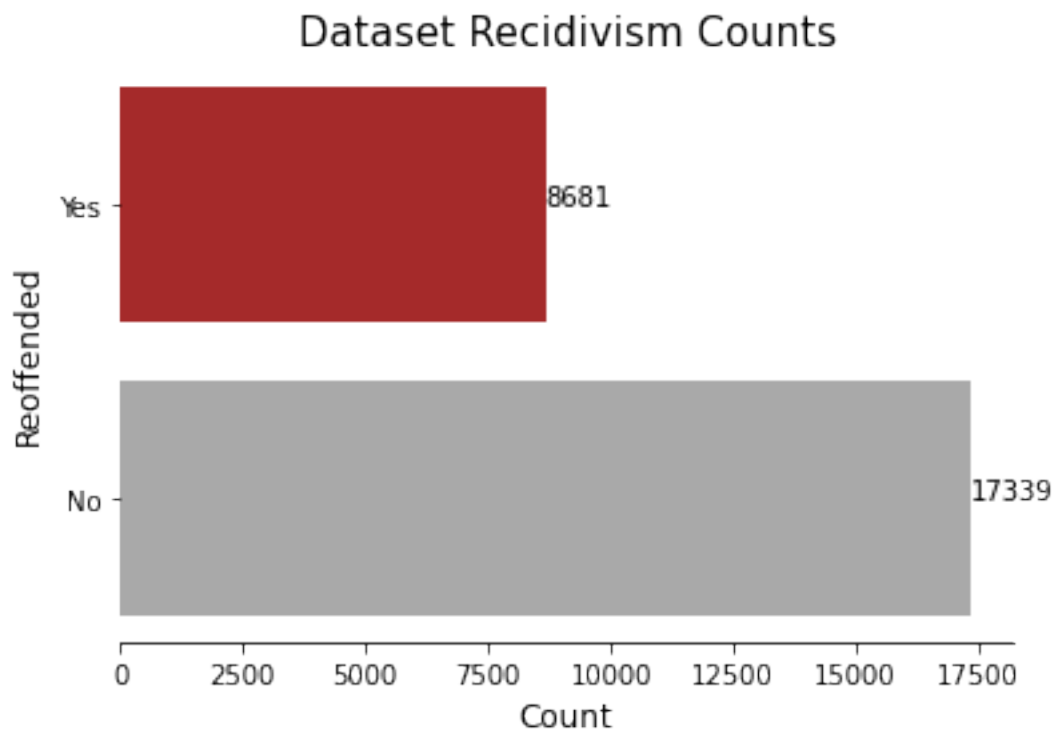
# Remove left border
left = ax.spines["left"]
left.set_visible(False)

# Plot
plt.barh(recid_counts.index, recid_counts['Return to Prison'],
         color=['darkgray', 'brown'])
plt.ylabel("Reoffended", size=12)
plt.xlabel("Count", size=12)
plt.title("Dataset Recidivism Counts", size=15)
plt.xticks([0, 1], ["No", "Yes"])

for index, value in enumerate(recid_counts['Return to Prison']):
    plt.text(value, index, str(value))

plt.show()

```



```
[3]: # Create new df dropping columns that are not going to be used
clean_df = df.drop(columns = ['Recidivism Reporting Year', 'Main Supervising_
    ↳District', 'Race - Ethnicity', 'Days to Return',
    'Recidivism Type', 'New Offense Classification', 'New_
    ↳Offense Type', 'New Offense Sub Type',
    'Target Population'])
clean_df.head()
```

```
[3]:
```

	Fiscal Year Released	Release Type	Age At Release	Sex	\
0	2010	Parole	25-34	Male	
1	2010	Discharged - End of Sentence	25-34	Male	
2	2010	Parole	35-44	Male	
3	2010	Parole	25-34	Male	
4	2010	Discharged - End of Sentence	35-44	Male	

	Offense Classification	Offense Type	Offense Subtype	Return to Prison	
0	C Felony	Violent	Robbery	Yes	
1	D Felony	Property	Theft	Yes	
2	B Felony	Drug	Trafficking	Yes	
3	B Felony	Other	Other Criminal	No	
4	D Felony	Violent	Assault	Yes	

```
[4]: # Change column names to be more usable
clean_df = clean_df.rename(columns={"Fiscal Year Released": "year", "Release_
    ↳Type": "release", "Age At Release ": "age",
    "Sex": "sex", "Offense Classification":_
    ↳"classification", "Offense Type": "type",
    "Offense Subtype": "subtype", "Return to_
    ↳Prison": "reoffend"})
clean_df.head()
```

```
[4]:
```

	year	release	age	sex	classification	type	\
0	2010	Parole	25-34	Male	C Felony	Violent	
1	2010	Discharged - End of Sentence	25-34	Male	D Felony	Property	
2	2010	Parole	35-44	Male	B Felony	Drug	
3	2010	Parole	25-34	Male	B Felony	Other	
4	2010	Discharged - End of Sentence	35-44	Male	D Felony	Violent	

	subtype	reoffend
0	Robbery	Yes
1	Theft	Yes
2	Trafficking	Yes
3	Other Criminal	No
4	Assault	Yes

```
[6]: # Look at unique values for columns
clean_df.release.unique()
```

```
[6]: array(['Parole', 'Discharged - End of Sentence', 'Special Sentence', nan,
        'Interstate Compact Parole', 'Parole Granted',
        'Discharged - Expiration of Sentence',
        'Paroled w/Immediate Discharge', 'Paroled to Detainer - Iowa',
        'Paroled to Detainer - U.S. Marshall',
        'Paroled to Detainer - Out of State',
        'Released to Special Sentence', 'Paroled to Detainer - INS'],
        dtype=object)
```

```
[7]: # Look at unique values for columns
clean_df.classification.unique()
```

```
[7]: array(['C Felony', 'D Felony', 'B Felony', 'Felony - Enhanced',
        'Aggravated Misdemeanor', 'Other Felony (Old Code)',
        'Serious Misdemeanor', 'Sexual Predator Community Supervision',
        'Simple Misdemeanor', 'Felony - Enhancement to Original Penalty',
        'Special Sentence 2005', 'Felony - Mandatory Minimum',
        'Other Felony', 'A Felony', 'Other Misdemeanor'], dtype=object)
```

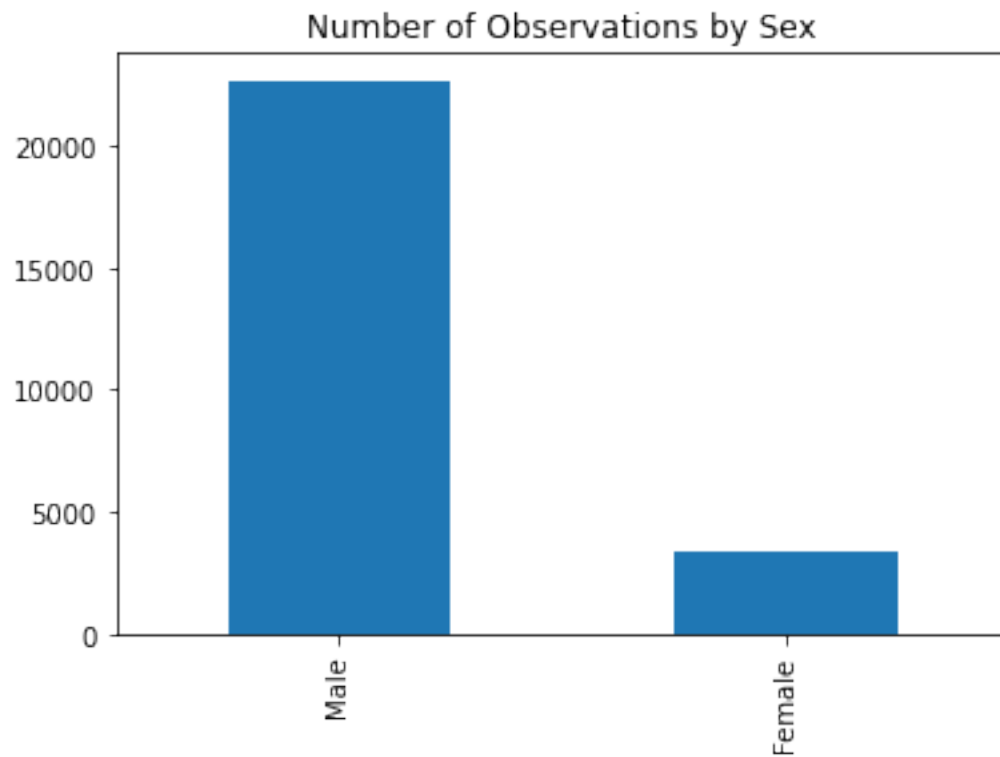
```
[5]: # The way records were kept changed over time. Combine values that are the same
      ↪but labeled differently.
clean_df2 = clean_df.replace({'release': {'Discharged - End of Sentence':
      ↪'Discharged - Expiration of Sentence',
      'Parole': 'Parole Granted', 'Special
      ↪Sentence': 'Released to Special Sentence'}})

clean_df2 = clean_df2.replace({'classification': {'Other Felony (Old Code)':
      ↪'Other Felony'}})

# Replace nan with Unknown
clean_df2['release'].fillna('Unknown', inplace=True)
```

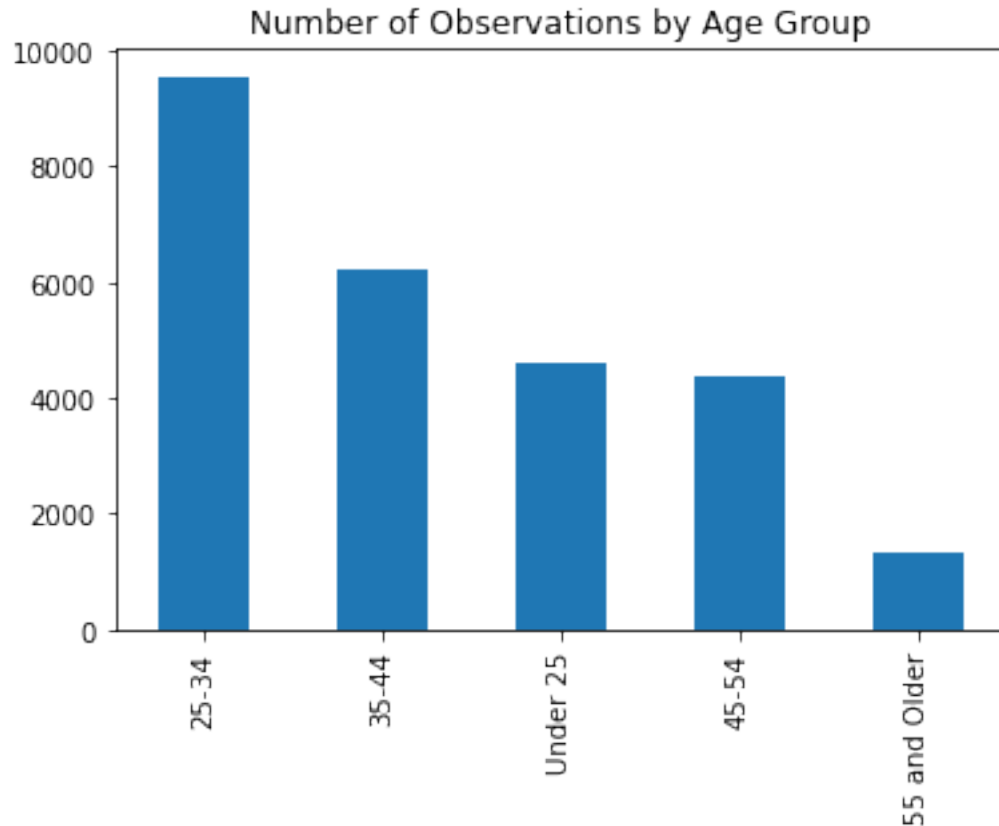
```
[6]: clean_df2.sex.value_counts().plot(kind="bar")
plt.title("Number of Observations by Sex")
```

```
[6]: Text(0.5, 1.0, 'Number of Observations by Sex')
```



```
[10]: clean_df2.age.value_counts().plot(kind="bar")  
      plt.title("Number of Observations by Age Group")
```

```
[10]: Text(0.5, 1.0, 'Number of Observations by Age Group')
```



```
[6]: # Look to see how many nans are left and replace them with mode values
clean_df2.isna().sum()
```

```
# Replace nan in sex with mode
clean_df2['sex'].fillna('Male', inplace=True)
```

```
# Replace nan in age with mode
clean_df2['age'].fillna('25-34', inplace=True)
```


```
[7]: # Create binary target vector indicating if class 0
clean_df2['reoffend'] = clean_df2['reoffend'].replace({"No": 0, "Yes": 1})
clean_df2.head()
```

```
[7]:   year      release  age  sex classification \
0  2010      Parole Granted  25-34  Male      C Felony
1  2010  Discharged - Expiration of Sentence  25-34  Male      D Felony
2  2010      Parole Granted  35-44  Male      B Felony
3  2010      Parole Granted  25-34  Male      B Felony
4  2010  Discharged - Expiration of Sentence  35-44  Male      D Felony
```

	type	subtype	reoffend
0	Violent	Robbery	1
1	Property	Theft	1
2	Drug	Trafficking	1
3	Other	Other Criminal	0
4	Violent	Assault	1

```
[12]: # Look at density plot of target variable using seaborn
import seaborn as sns
sns.distplot(clean_df2.reoffend)
```

```
[12]: <AxesSubplot:xlabel='reoffend', ylabel='Density'>
```



```
[8]: # Write cleaned up dataframe to excel file
clean_df2.to_excel("recidivism_cleaned.xlsx")
```

```
[8]: # Change categorical features to numerical values
#categorical data
categorical_cols = ['release', 'age', 'sex', 'classification', 'type', '
↳ 'subtype']

# Use get dummies to one-hot encode the categorical columns
clean_df3 = pd.get_dummies(clean_df2, columns = categorical_cols)
clean_df3
```



```

[8]:      year  reoffend  release_Discharged - Expiration of Sentence  \
0      2010          1                                0
1      2010          1                                1
2      2010          1                                0
3      2010          0                                0
4      2010          1                                1
...
26015  2015          0                                0
26016  2015          0                                0
26017  2015          0                                0
26018  2015          0                                0
26019  2015          1                                0

```

```

      release_Interstate Compact Parole  release_Parole Granted  \
0                                0                                1
1                                0                                0
2                                0                                1
3                                0                                1
4                                0                                0
...
26015                                0                                0
26016                                0                                0
26017                                0                                1
26018                                0                                0
26019                                0                                0

```

```

      release_Paroled to Detainer - INS  release_Paroled to Detainer - Iowa  \
0                                0                                0
1                                0                                0
2                                0                                0
3                                0                                0
4                                0                                0
...
26015                                1                                0
26016                                0                                0
26017                                0                                0
26018                                0                                0
26019                                0                                0

```

```

      release_Paroled to Detainer - Out of State  \
0                                0
1                                0
2                                0
3                                0
4                                0
...
26015                                0

```

26016	0
26017	0
26018	0
26019	1

	release_Paroled to Detainer - U.S. Marshall \
0	0
1	0
2	0
3	0
4	0
...	...
26015	0
26016	0
26017	0
26018	0
26019	0

	release_Paroled w/Immediate Discharge	...	subtype_Robbery \
0	0	...	1
1	0	...	0
2	0	...	0
3	0	...	0
4	0	...	0
...	...	...	...
26015	0	...	0
26016	0	...	0
26017	0	...	0
26018	1	...	0
26019	0	...	0

	subtype_Sex	subtype_Sex Offender Registry/Residency \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...	...	...
26015	0	0
26016	1	0
26017	0	0
26018	0	0
26019	0	0

	subtype_Special Sentence Revocation	subtype_Stolen Property \
0	0	0
1	0	0

2	0	0
3	0	0
4	0	0
...	...	...
26015	0	0
26016	0	0
26017	0	0
26018	0	0
26019	0	0

	subtype_Theft	subtype_Traffic	subtype_Trafficking	subtype_Vandalism	\
0	0	0	0	0	
1	1	0	0	0	
2	0	0	1	0	
3	0	0	0	0	
4	0	0	0	0	
...	...	...	...	...	
26015	0	0	0	0	
26016	0	0	0	0	
26017	0	1	0	0	
26018	1	0	0	0	
26019	0	0	0	0	

	subtype_Weapons
0	0
1	0
2	0
3	0
4	0
...	...
26015	0
26016	0
26017	0
26018	0
26019	0

[26020 rows x 64 columns]

```
[9]: # Move the reoffend column to the last column
clean_df4 = clean_df3[[c for c in clean_df3 if c not in ['reoffend']]
               + ['reoffend']]
clean_df4
```

```
[9]:      year  release_Discharged - Expiration of Sentence  \
0      2010                                           0
1      2010                                           1
2      2010                                           0
```

3	2010	0
4	2010	1
...	...	...
26015	2015	0
26016	2015	0
26017	2015	0
26018	2015	0
26019	2015	0

	release_Interstate Compact Parole	release_Parole Granted \
0	0	1
1	0	0
2	0	1
3	0	1
4	0	0
...	...	...
26015	0	0
26016	0	0
26017	0	1
26018	0	0
26019	0	0

	release_Paroled to Detainer - INS	release_Paroled to Detainer - Iowa \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...	...	...
26015	1	0
26016	0	0
26017	0	0
26018	0	0
26019	0	0

	release_Paroled to Detainer - Out of State \
0	0
1	0
2	0
3	0
4	0
...	...
26015	0
26016	0
26017	0
26018	0
26019	1

	release_Paroled to Detainer - U.S. Marshall \
0	0
1	0
2	0
3	0
4	0
...	...
26015	0
26016	0
26017	0
26018	0
26019	0

	release_Paroled w/Immediate Discharge \
0	0
1	0
2	0
3	0
4	0
...	...
26015	0
26016	0
26017	0
26018	1
26019	0

	release_Released to Special Sentence ...	subtype_Sex \
0	0 ...	0
1	0 ...	0
2	0 ...	0
3	0 ...	0
4	0 ...	0
...	... ...	...
26015	0 ...	0
26016	1 ...	1
26017	0 ...	0
26018	0 ...	0
26019	0 ...	0

	subtype_Sex Offender Registry/Residency \
0	0
1	0
2	0
3	0
4	0
...	...

26015	0
26016	0
26017	0
26018	0
26019	0

	subtype_Special Sentence Revocation	subtype_Stolen Property \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...	...	...
26015	0	0
26016	0	0
26017	0	0
26018	0	0
26019	0	0

	subtype_Theft	subtype_Traffic	subtype_Trafficking	subtype_Vandalism \
0	0	0	0	0
1	1	0	0	0
2	0	0	1	0
3	0	0	0	0
4	0	0	0	0
...	...	...	...	...
26015	0	0	0	0
26016	0	0	0	0
26017	0	1	0	0
26018	1	0	0	0
26019	0	0	0	0

	subtype_Weapons	reoffend
0	0	1
1	0	1
2	0	1
3	0	0
4	0	1
...	...	...
26015	0	0
26016	0	0
26017	0	0
26018	0	0
26019	0	1

[26020 rows x 64 columns]

```
[10]: # Split data into features and target data
features = clean_df4.iloc[:, :-1]
target = clean_df4.iloc[:, -1:]

[11]: # Use RandomForestClassifier for feature selection because the dataset contains
      ↪ both numeric and categorical features
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_selection import SelectFromModel

# Specify random forest instance, indicating the number of trees
sel = SelectFromModel(RandomForestClassifier(n_estimators = 100))

# Use selectFromModel object to automatically select features
sel.fit(features, target.values.ravel())

# Make a list and count the selected features
selected_feat= features.columns[(sel.get_support())]

print("Best Number of Features: ", len(selected_feat))
print("List of Feature Names: ", selected_feat)
```

Best Number of Features: 16

List of Feature Names: Index(['year', 'release\_Discharged - Expiration of Sentence',  
'release\_Parole Granted', 'release\_Released to Special Sentence',  
'release\_Unknown', 'age\_25-34', 'age\_35-44', 'age\_45-54',  
'age\_55 and Older', 'age\_Under 25', 'sex\_Female', 'sex\_Male',  
'classification\_Aggravated Misdemeanor', 'classification\_C Felony',  
'classification\_D Felony', 'type\_Violent'],  
dtype='object')

```
[12]: # Create a features dataframe that has the identified best features
best_features = features[['year', 'release_Discharged - Expiration of Sentence',  

'release_Parole Granted', 'release_Released to Special Sentence',  

'release_Unknown', 'age_25-34', 'age_35-44', 'age_45-54',  

'age_55 and Older', 'age_Under 25', 'sex_Female', 'sex_Male',  

'classification_Aggravated Misdemeanor', 'classification_C Felony',  

'classification_D Felony', 'type_Violent']]
best_features
```

```
[12]:
```

	year	release_Discharged - Expiration of Sentence	\
0	2010		0
1	2010		1
2	2010		0
3	2010		0
4	2010		1
...	...		...

26015	2015	0
26016	2015	0
26017	2015	0
26018	2015	0
26019	2015	0

	release_Parole Granted	release_Released to Special Sentence	\
0	1	0	
1	0	0	
2	1	0	
3	1	0	
4	0	0	
...	...	...	
26015	0	0	
26016	0	1	
26017	1	0	
26018	0	0	
26019	0	0	

	release_Unknown	age_25-34	age_35-44	age_45-54	age_55 and Older	\
0	0	1	0	0	0	
1	0	1	0	0	0	
2	0	0	1	0	0	
3	0	1	0	0	0	
4	0	0	1	0	0	
...	...	...	...	...	...	
26015	0	0	0	0	0	
26016	0	0	1	0	0	
26017	0	1	0	0	0	
26018	0	1	0	0	0	
26019	0	0	1	0	0	

	age_Under 25	sex_Female	sex_Male	\
0	0	0	1	
1	0	0	1	
2	0	0	1	
3	0	0	1	
4	0	0	1	
...	...	...	...	
26015	1	0	1	
26016	0	0	1	
26017	0	1	0	
26018	0	0	1	
26019	0	0	1	

	classification_Aggravated Misdemeanor	classification_C Felony	\
0	0	1	



1	0	0
2	0	0
3	0	0
4	0	0
...	...	...
26015	0	1
26016	0	1
26017	1	0
26018	0	0
26019	0	0

	classification_D Felony	type_Violent
0	0	1
1	1	0
2	0	0
3	0	0
4	1	1
...	...	...
26015	0	1
26016	0	1
26017	0	0
26018	1	0
26019	1	1

[26020 rows x 16 columns]

```
[13]: # reset index
best_features.reset_index(inplace=True, drop=True)
```

```
[14]: # Train Test Split
from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(best_features, target,
↳test_size =0.3, random_state=11)
```

```
[23]: # number of samples in each set
print("No. of samples in training set: ", X_train.shape[0])
print("No. of samples in validation set:", X_val.shape[0])
```

No. of samples in training set: 18214  
No. of samples in validation set: 7806

```
[24]: # Check if training target vector had imbalanced classes
y_train['reoffend'].value_counts()
```

```
[24]: 0    12186
      1     6028
```

Name: reoffend, dtype: int64

### *Upsampling Training Data*

```
[15]: # To fix imbalanced classes, upsample training data.
from sklearn.utils import resample

#combine them back for resampling
train_data = pd.concat([X_train, y_train], axis=1)

# separate minority and majority classes
negative = train_data[train_data.reoffend==0]
positive = train_data[train_data.reoffend==1]

# upsample minority
pos_upsampled = resample(positive,
    replace=True, # sample with replacement
    n_samples=len(negative), # match number in majority class
    random_state=27) # reproducible results

# combine majority and upsampled minority
upsampled = pd.concat([negative, pos_upsampled])

# check new class counts
upsampled.reoffend.value_counts()
```

```
[15]: 0    12186
      1    12186
      Name: reoffend, dtype: int64
```

```
[16]: # Reassign upsampled training features and training target
X_train = upsampled.drop("reoffend",axis=1)
y_train = upsampled["reoffend"]
```

```
[17]: # number of samples in each set
print("No. of samples in training set: ", X_train.shape[0])
print("No. of samples in validation set:", X_val.shape[0])
```

No. of samples in training set: 24372  
No. of samples in validation set: 7806

### *Correlation Matrix of Best Features*

```
[49]: # Compute the correlation matrix
corr = train_data.corr()

# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))
```

```


# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

cmap = sns.light_palette("Brown", reverse=False, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})

```

[49]: <AxesSubplot:>




```

[50]: # What I'm interested is any correlation the attributes might have with the
      ↪ Class
      # Compute the correlation matrix
      corr = train_data.corr()

```

```
plt.figure(figsize=(8, 12))
heatmap = sns.heatmap(train_data.corr()[['reoffend']].
    ↳sort_values(by='reoffend', ascending=False), vmin=0, vmax=1, annot=True,
    ↳cmap=cmap)
heatmap.set_title('Features Correlating with Class', fontdict={'fontsize':18},
    ↳pad=16);
```



[ ]:

### *Baseline Classification Model*

```
[19]: # Create a baseline classification model to use for comparison
from sklearn.dummy import DummyClassifier

# Create classifier
dummy = DummyClassifier(strategy='uniform', random_state=0)

# Train model
dummy.fit(X_train, y_train.values.ravel())

# Get accuracy score
dummy.score(X_val, y_val.values.ravel())
```

[19]: 0.49743786830643094

### *Random Forest Model*

```
[23]: from sklearn.ensemble import RandomForestClassifier

# Create random forest classifier object
randomforest = RandomForestClassifier(random_state=0, n_jobs=-1,
    ↪class_weight="balanced")

#Train model
rf_model = randomforest.fit(X_train, y_train.values.ravel())
```

```
[24]: # view model accuracy score
rf_model.score(X_val, y_val.values.ravel())
```

[24]: 0.5807071483474251

### *Hyperparameter Tuning*

```
[34]: # Use RandomizedSearchCV to tune hyperparameters because it is computationally
    ↪cheaper than an exhaustive search.
from sklearn.model_selection import RandomizedSearchCV

# Create range for number of trees in random forest
n_estimators = [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000]

# Create range for number of features to consider at every split
max_features = ['auto', 'sqrt']

# Create range for maximum number of levels in tree
max_depth = [10, 20, 30, 40, 50, 60, 70, 80, 90, 100, None]

# Create range for minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
```

```

# Create range for minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]

# Define method of selecting samples for training each tree
bootstrap = [True, False]

# Create hyperparameter options
hyperparameters = dict(n_estimators=n_estimators, max_features=max_features,
    ↪max_depth=max_depth,
                        min_samples_split=min_samples_split,
    ↪min_samples_leaf=min_samples_leaf, bootstrap=bootstrap)

# Create randomized search
randomizedsearch = RandomizedSearchCV(rf_model, hyperparameters,
    ↪random_state=1, n_iter=100, cv=5, verbose=0, n_jobs=-1)

# Fit randomized search
best_rf_model = randomizedsearch.fit(X_train, y_train.values.ravel())

```

```

[35]: # View best hyperparameters
best_rf_model.best_params_

```

```

[35]: {'n_estimators': 400,
      'min_samples_split': 5,
      'min_samples_leaf': 4,
      'max_features': 'sqrt',
      'max_depth': 10,
      'bootstrap': True}

```

```

[36]: # Get accuracy score for best_model and compare to intial model
best_rf_model.score(X_val, y_val.values.ravel())

```

```

[36]: 0.6026133743274404

```

### Adaboost

```

[20]: # Try to use boosting in attempt to improve the model
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import cross_val_score

# Create adaboost tree classifier object
adaboost = AdaBoostClassifier(random_state=0)

# Train model on the important features data set
ab_model = adaboost.fit(X_train, y_train.values.ravel())

# Cross-validate Adaboost model

```

```

cvScore = cross_val_score(ab_model, X_val, y_val.values.ravel(),
    ↪scoring="accuracy")
print("CV Scores:", cvScore)
print("CV Scores Mean:", cvScore.mean())

```

CV Scores: [0.66133163 0.66495836 0.66367713 0.65406791 0.66303652]  
 CV Scores Mean: 0.661414307286852

```

[30]: # Tune hyperparameters with Adaboost
from sklearn.model_selection import GridSearchCV

# Define hyperparameter candidates
param_grid = {'n_estimators': [100, 200, 400, 600, 800, 1000],
              'learning_rate': [0.001, 0.01, 0.1, 0.2, 0.5]}

# Instantiate GridSearchCV model
gs_ab = GridSearchCV(AdaBoostClassifier(), param_grid = param_grid)

# Train model
gs_ab_model = gs_ab.fit(X_train, y_train.values.ravel())

# Cross-validate model
cvScore2 = cross_val_score(gs_ab_model, X_val, y_val.values.ravel(),
    ↪scoring="accuracy")
print("CV Scores:", cvScore2)
print("CV Scores Mean:", cvScore2.mean())

```

CV Scores: [0.66517286 0.66495836 0.66303652 0.65150545 0.66303652]  
 CV Scores Mean: 0.6615419381351295

```

[32]: # View best hyperparameters
gs_ab_model.best_params_

```

```

[32]: {'learning_rate': 0.1, 'n_estimators': 600}

```

### *Skip to Here to Run Final Model & Results*

```

[18]: ### After upsampling & changing target factoring run Adaboost with best params
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import cross_val_score

# Create classifier object
myBoost = AdaBoostClassifier(random_state=0, n_estimators=600, learning_rate=0.
    ↪1)

# Train model
final_model = myBoost.fit(X_train, y_train.values.ravel())

```

```

# Cross-validate model
cvScore2 = cross_val_score(final_model, X_val, y_val.values.ravel(),
    ↪scoring="accuracy")
print("CV Scores:", cvScore2)
print("CV Scores Mean:", cvScore2.mean())

```

CV Scores: [0.66389245 0.66559898 0.66367713 0.65663037 0.6623959 ]  
 CV Scores Mean: 0.6624389631716102

```

[19]: # Compute predictions
from sklearn.metrics import roc_auc_score, roc_curve, auc, plot_roc_curve,\
    precision_score, recall_score, f1_score, classification_report,
    ↪precision_recall_curve, plot_precision_recall_curve, \
    average_precision_score

# compute predict probabilities on test set
# only need the probabilities for the positive class only
y_pred_probs = final_model.predict_proba(X_val)[:, 1]

# compute predictions on test set
y_preds = final_model.predict(X_val)

```

```

[20]: # Compute AUC
auc_score = roc_auc_score(y_val, y_pred_probs)
round(auc_score,4)

```


[20]: 0.6427

```

[22]: # Plot ROC Curve
lw = 2
plot_roc_curve(final_model, X_val, y_val, name = 'AdaBoost',
    color="red", pos_label=None)
plt.plot([0, 1], [0, 1], 'k--', lw=lw)
plt.title('AdaBoost ROC Curve')
plt.show()


```





```
[23]: # Plot predictions on the test set
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay

# Create confusion matrix
cm = confusion_matrix(y_val, y_preds)
cm_display = ConfusionMatrixDisplay(cm).plot()
```




```
[24]: # Build Classification report
print(classification_report(y_val, y_preds))
```

	precision	recall	f1-score	support
0	0.75	0.59	0.66	5153
1	0.44	0.62	0.51	2653
accuracy			0.60	7806
macro avg	0.59	0.60	0.59	7806
weighted avg	0.64	0.60	0.61	7806

```
[26]: # Compute Average Precision Score
y_score = final_model.decision_function(X_val)

average_precision = average_precision_score(y_val, y_score)

# Plot PR curve
disp = plot_precision_recall_curve(final_model, X_val, y_val, name=None,
    pos_label=1, color="red")
disp.ax_.set_title('AdaBoost Precision-Recall Curve')
plt.show()
```



### *Additional Visualizations for Presentation Below*

```
[49]: plt.rcParams['figure.figsize'] = (30, 10)

# make subplots
fig, axes = plt.subplots(nrows = 1, ncols = 2)

# make the data for Classification to feed into the visulizer
Class_No = clean_df2.replace({'reoffend': {1: 'reoffend', 0: 'Yes'}})[clean_df2['reoffend']==1]['classification'].value_counts()
Class_Yes = clean_df2.replace({'reoffend': {1: 'reoffend', 0: 'No'}})[clean_df2['reoffend']==0]['classification'].value_counts()
Class_Yes = Class_Yes.reindex(index = Class_No.index)
# make the bar plot
p5 = axes[0].bar(Class_No.index, Class_No.values, color='darkgray')
p6 = axes[0].bar(Class_Yes.index, Class_Yes.values, bottom=Class_No.values, color='brown')
axes[0].set_title('Classification', fontsize=25)
axes[0].set_ylabel('Counts', fontsize=20)
axes[0].tick_params(axis='y', labels=15)
axes[0].tick_params(axis='x', labels=15, rotation=90)
axes[0].legend((p5[0], p6[0]), ('Did Not Reoffend', 'Reoffended'), fontsize = 15, loc='upper right')
```

```

axes[0].bar_label(p6, label_type='center', color='white', weight="bold")


# make the data for Release Type to feed into the visulizer
Release_No = clean_df2.replace({'reoffend': {1: 'reoffend', 0: 'Yes'}})[clean_df2['reoffend']==1]['release'].value_counts()
Release_Yes = clean_df2.replace({'reoffend': {1: 'reoffend', 0: 'No'}})[clean_df2['reoffend']==0]['release'].value_counts()
Release_Yes = Release_Yes.reindex(index = Release_No.index)
# make the bar plot
p9 = axes[1].bar(Release_No.index, Release_No.values, color='darkgray')
p10 = axes[1].bar(Release_Yes.index, Release_Yes.values, color='brown')
axes[1].set_title('Release Type', fontsize=25)
axes[1].set_ylabel('Counts', fontsize=20)
axes[1].tick_params(axis='y', labels=15)
axes[1].tick_params(axis='x', labels=15, rotation=90)
axes[1].legend((p9[0], p10[0]), ('Did Not Reoffend', 'Reoffended'), fontsize = 15, loc='upper right')
axes[1].bar_label(p10, label_type='center', color='white', weight="bold")

```

```

[49]: [Text(0, 0, '5914'),
Text(0, 0, '1866'),
Text(0, 0, '350'),
Text(0, 0, '388'),
Text(0, 0, '73'),
Text(0, 0, '4'),
Text(0, 0, '43'),
Text(0, 0, '5'),
Text(0, 0, '38'),
Text(0, 0, 'nan')]

```



```

[51]: # Split the stacked bar charts apart
#set up the figure size
plt.rcParams['figure.figsize'] = (30, 20)

# make subplots
fig, axes = plt.subplots(nrows = 2, ncols = 2)

# make the data for Reoffend to feed into the visulizer
axes[0, 0].barh(recid_counts.index, recid_counts['Return to Prison'],
    ↪color=['darkgray', 'brown'])
axes[0, 0].set_ylabel("Reoffended", size=20)
axes[0, 0].set_xlabel("Count", size=20)
axes[0, 0].set_title("Recidivism Counts", size=25, weight="bold")
axes[0, 0].set_yticklabels((" ", " ", "No", " ", " ", " ", "Yes"))
axes[0, 0].set_xlim(0, 19000)
for index, value in enumerate(recid_counts['Return to Prison']):
    axes[0, 0].text(value, index, str(value), weight="bold")

# make the data for Gender to feed into the visulizer
Sex_No = clean_df2.replace({'reoffend': {1: 'reoffend', 0:
    ↪'Yes'}})[clean_df2['reoffend']==1]['sex'].value_counts()
Sex_Yes = clean_df2.replace({'reoffend': {1: 'reoffend', 0:
    ↪'No'}})[clean_df2['reoffend']==0]['sex'].value_counts()
Sex_Yes = Sex_Yes.reindex(index = Sex_No.index)
# make the bar plot
p3 = axes[0, 1].bar(Sex_No.index, Sex_No.values, color='darkgray')
p4 = axes[0, 1].bar(Sex_Yes.index, Sex_Yes.values, bottom=Sex_No.values,
    ↪color='brown')
axes[0, 1].set_title('Gender', fontsize=25, weight="bold")
axes[0, 1].set_ylabel('Counts', fontsize=20)
axes[0, 1].tick_params(axis='both', labelsize=15)
axes[0, 1].legend((p3[0], p4[0]), ('Did Not Reoffend', 'Reoffended'), fontsize
    ↪= 15)
axes[0, 1].bar_label(p4, label_type='center', color='white', weight="bold")

# make the data for Age to feed into the visulizer
Age_No = clean_df2.replace({'reoffend': {1: 'reoffend', 0:
    ↪'Yes'}})[clean_df2['reoffend']==1]['age'].value_counts()
Age_Yes = clean_df2.replace({'reoffend': {1: 'reoffend', 0:
    ↪'No'}})[clean_df2['reoffend']==0]['age'].value_counts()
Age_Yes = Age_Yes.reindex(index = Age_No.index)
# make the bar plot
p1 = axes[1, 0].bar(Age_No.index, Age_No.values, color='darkgray')
p2 = axes[1, 0].bar(Age_Yes.index, Age_Yes.values, bottom=Age_No.values,
    ↪color='brown')

```

```

axes[1, 0].set_title('Age', fontsize=25, weight="bold")
axes[1, 0].set_ylabel('Counts', fontsize=20)
axes[1, 0].tick_params(axis='both', labelsize=15)
axes[1, 0].legend((p1[0], p2[0]), ('Did Not Reoffend', 'Reoffended'), fontsize=
    ↪ 15)
axes[1, 0].bar_label(p2, label_type='center', color='white', weight="bold")

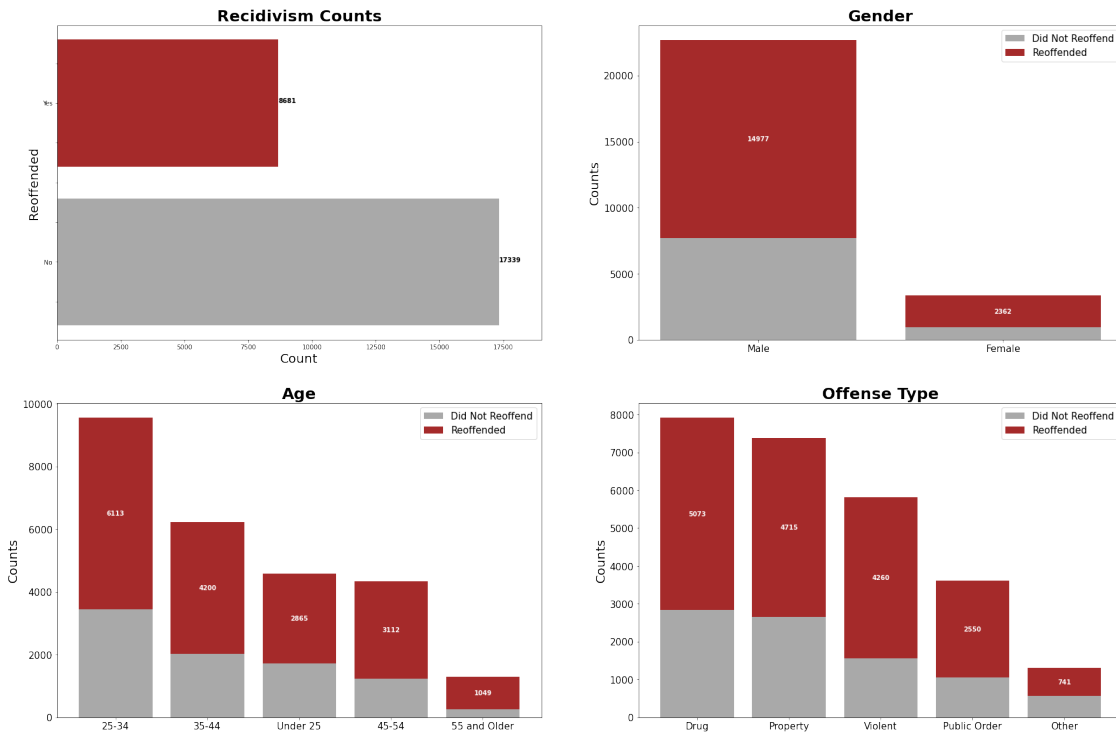
# make the data for Offense Type to feed into the visulizer
Offense_No = clean_df2.replace({'reoffend': {1: 'reoffend', 0:
    ↪ 'Yes'}})[clean_df2['reoffend']==1]['type'].value_counts()
Offense_Yes = clean_df2.replace({'reoffend': {1: 'reoffend', 0:
    ↪ 'No'}})[clean_df2['reoffend']==0]['type'].value_counts()
Offense_Yes = Offense_Yes.reindex(index = Offense_No.index)
# make the bar plot
p7 = axes[1, 1].bar(Offense_No.index, Offense_No.values, color='darkgray')
p8 = axes[1, 1].bar(Offense_Yes.index, Offense_Yes.values, bottom=Offense_No.
    ↪ values, color='brown')
axes[1, 1].set_title('Offense Type', fontsize=25, weight="bold")
axes[1, 1].set_ylabel('Counts', fontsize=20)
axes[1, 1].tick_params(axis='both', labelsize=15)
axes[1, 1].legend((p7[0], p8[0]), ('Did Not Reoffend', 'Reoffended'), fontsize=
    ↪ 15)
axes[1, 1].bar_label(p8, label_type='center', color='white', weight="bold")

plt.show()

```

<ipython-input-51-ed512fe37307>:13: UserWarning: FixedFormatter should only be used together with FixedLocator

```
axes[0, 0].set_yticklabels(("", "", "No", "", "", "", "Yes"))
```



```
[56]: # Create table with Pearson Chi-square and Cramers V between best features and
      ↪ target variable. (From Christine).
import researchpy as rp

# Create new DF to store results
new_test = pd.DataFrame(columns=['feature', 'Pearson Chi-square', \
                                'p-value', 'Cramers V', \
                                'Cramer_val'])

# Define columns to use
cols = ['release', 'year', 'age', 'sex', 'classification', 'type', 'subtype']

# Iterate through columns and calculate statistics
for i in cols:
    crosstab, test_results, expected = rp.crosstab(
        clean_df2[i], clean_df2['reoffend'],
        test= "chi-square",
        expected_freqs= True,
        prop= "cell")

    # Create new var to append to new test df
    if test_results.iloc[2,1] > .25:
        new_var = 'Very Strong'
    elif test_results.iloc[2,1] > .15:
        new_var = 'Strong'
```

```

elif test_results.iloc[2,1] > .10:
    new_var = 'Moderate'
elif test_results.iloc[2,1] > .05:
    new_var = 'Weak'
else:
    new_var = 'No or very weak'

# Append row to new_test DF
new_row = {'feature':i, \
           'Pearson Chi-square': test_results.iloc[0,1], \
           'p-value': test_results.iloc[1,1], \
           'Cramers V': test_results.iloc[2,1], \
           'Cramer_val': new_var}

new_test = new_test.append(new_row, ignore_index=True)

# https://www.pythonfordatascience.org/chi-square-test-of-independence-python/
# Phi and Cramer's V Interpretation
# >0.25 Very strong
# >0.15 Strong
# >0.10 Moderate
# >0.05 Weak
# >0 No or very weak

# View results
new_test

```

```

[56]:

```

	feature	Pearson Chi-square	p-value	Cramers V	Cramer_val
0	release	771.7566	0.0	0.1722	Strong
1	year	93.4152	0.0	0.0599	Weak
2	age	230.2147	0.0	0.0941	Weak
3	sex	28.9993	0.0	0.0334	No or very weak
4	classification	90.9506	0.0	0.0591	Weak
5	type	247.1988	0.0	0.0975	Weak
6	subtype	355.6772	0.0	0.1169	Moderate

```

[27]: # Compute correlation coefficients using the pearson method and create a mask for
# input to seaborn to plot one-half of the heatmap
corrmat = best_features.corr(method='pearson')

ab_mask = np.triu(np.ones(corrmat.shape)).astype(np.bool)
title = "AdaBoost Best Features Correlation Heatmap"

plt.rcParams['font.size'] = (16)
with sns.axes_style("white"):
    f, ax = plt.subplots(figsize=(20, 10))
    ax = sns.heatmap(corrmat,

```



```

        mask=ab_mask,
        cmap="RdGy",
        annot=True,
        cbar=None,
    )

ax.set_xticklabels(
    ax.get_xticklabels(),
    rotation=25,
    horizontalalignment='right')
ax.set_title(title, fontsize=18, weight="bold" )

plt.show()

```

<ipython-input-27-2f90113f0912>:5: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool\_` here.

Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
ab_mask = np.triu(np.ones(corrmat.shape)).astype(np.bool)
```

