

CKD_Predict_notebook

February 10, 2023

```
[203]: # Project 2 CKD Prediction
# Christine Drosco 1/30/2023
'''
Predict the probabiity of a person being diagnosed with CKD
'''
```

```
[203]: '\nPredict the probabiity of a person being diagnosed with CKD\n\n'
```

```
[8]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import numpy as np
```

1 Perform Data Cleaning

```
[9]: # Read in data set
```

```
[72]: # The UCI dataset

df = pd.read_csv('ckd_dataset.csv', header=0)
```

```
[73]: df.head()
```

```
[73]:   age  bp    sg  al  su    rbc    pc    ba  bgr  bu  ...  wbcc  rbcc  \
0   48  80  1.020  1  0      ?   normal  notpresent  121  36  ...  7800  5.2
1    7  50  1.020  4  0      ?   normal  notpresent    ?  18  ...  6000    ?
2   62  80  1.010  2  3  normal   normal  notpresent  423  53  ...  7500    ?
3   48  70  1.005  4  0  normal  abnormal  notpresent  117  56  ...  6700  3.9
4   51  80  1.010  2  0  normal   normal  notpresent  106  26  ...  7300  4.6
```

```
      htn  dm  cad  appet  pe  ane  class  Unnamed: 24
0  yes  yes  no  good  no  no  ckd      NaN
1   no   no  no  good  no  no  ckd      NaN
2   no  yes  no  poor  no  yes  ckd      NaN
```

```

3  yes  no  no  poor  yes  yes  ckd      NaN
4  no   no  no  good  no   no  ckd      NaN

```

[5 rows x 25 columns]

```
[74]: #Define column names with abbreviations
```

```

labels = {
    'age': 'age',
    'bp': 'blood_pressure',
    'sg': 'specific_gravity',
    'al': 'albumin',
    'su': 'sugar',
    'rbc': 'red_blood_cells',
    'pc': 'pus_cell',
    'pcc': 'pus_cell_clumps',
    'ba': 'bacteria',
    'bgr': 'blood_glucose_random',
    'bu': 'blood_urea',
    'sc': 'serum_creatinine',
    'sod': 'sodium',
    'pot': 'potassium',
    'hemo': 'hemoglobin',
    'pcv': 'packed_cell_volume',
    'wbcc': 'white_blood_cell_count',
    'rbcc': 'red_blood_cell_count',
    'htn': 'hypertension',
    'dm': 'diabetes_mellitus',
    'cad': 'coronary_artery_disease',
    'appet': 'appetite',
    'pe': 'pedal_edema',
    'ane': 'anemia',
    'class': 'class'}

```

```
[75]: #Clean the dataset
```

```

# Remove bogus column Unnamed: 24
#df.drop(['Unnamed: 24'], axis=1, inplace=True)
df

```

```

[75]:   age  bp    sg  al  su    rbc    pc    ba  bgr  bu  ...  wbcc  \
0   48  80  1.020  1  0      ?   normal  notpresent  121  36  ...  7800
1    7  50  1.020  4  0      ?   normal  notpresent    ?  18  ...  6000
2   62  80  1.010  2  3  normal   normal  notpresent  423  53  ...  7500
3   48  70  1.005  4  0  normal  abnormal  notpresent  117  56  ...  6700
4   51  80  1.010  2  0  normal   normal  notpresent  106  26  ...  7300
..   ..  ..   ...  ..  ..   ...   ...   ...  ...  ...  ...  ...

```

```

395  55  80  1.020  0  0  normal    normal  notpresent  140  49  ...  6700
396  42  70  1.025  0  0  normal    normal  notpresent   75  31  ...  7800
397  12  80  1.020  0  0  normal    normal  notpresent  100  26  ...  6600
398  17  60  1.025  0  0  normal    normal  notpresent  114  50  ...  7200
399  58  80  1.025  0  0  normal    normal  notpresent  131  18  ...  6800

```

```

      rbcc  htn   dm cad appet   pe  ane   class Unnamed: 24
0      5.2  yes  yes  no   good  no   no    ckd          NaN
1      ?   no   no   no   good  no   no    ckd          NaN
2      ?   no  yes  no   poor  no  yes    ckd          NaN
3     3.9  yes  no   no   poor  yes  yes    ckd          NaN
4     4.6   no  no   no   good  no   no    ckd          NaN
..  ...  ...  ...  ..  ...  ...  ...  ...
395  4.9   no  no   no   good  no   no  notckd          NaN
396  6.2   no  no   no   good  no   no  notckd          NaN
397  5.4   no  no   no   good  no   no  notckd          NaN
398  5.9   no  no   no   good  no   no  notckd          NaN
399  6.1   no  no   no   good  no   no  notckd          NaN

```

[400 rows x 25 columns]

```
[76]: #Check the col defs
df.dtypes
```

```

[76]: age          object
bp            object
sg            object
al            object
su            object
rbc           object
pc            object
ba            object
bgr           object
bu            object
sc            object
sod           object
pot           object
hemo          object
pcv           object
wbcc          object
rbcc          object
htn           object
dm            object
cad           object
appet         object
pe            object
ane           object

```

```
class          object
Unnamed: 24    object
dtype: object
```

```
[77]: # Change unknown values with ?
      # Take the average or use the lowest common values.
      # First find any NAN values

      df.isnull().sum()
```

```
[77]: age          0
      bp          0
      sg          0
      al          0
      su          0
      rbc         0
      pc          0
      ba          0
      bgr         0
      bu          0
      sc          0
      sod         0
      pot         0
      hemo        0
      pcv         0
      wbcc        0
      rbcc        0
      htn         0
      dm          1
      cad         0
      appet       0
      pe          0
      ane         0
      class       0
      Unnamed: 24  399
      dtype: int64
```

```
[78]: # Look at the occurrence of the values
      df['rbcc'].describe()
```

```
[78]: count      400
      unique     50
      top        ?
      freq      130
      Name: rbcc, dtype: object
```

```
[79]: df['su']
```

```
[79]: 0      0
      1      0
      2      3
      3      0
      4      0
      ..
      395    0
      396    0
      397    0
      398    0
      399    0
      Name: su, Length: 400, dtype: object
```

```
[80]: #Fill NaN value
      df['dm'].fillna('normal',inplace=True)
```

```
[81]: # Replace the Question marks.
      df['rbc'] = df['rbc'].replace({"?":"normal"})
```

```
[82]: # Replace bgr
      df['bgr'] = df['bgr'].replace({"?":150})
```

```
[83]: # Replace rbcc
      df['rbcc'] = df['rbcc'].replace({"?":5.0})
```

```
[84]: # Look for other columns with ? mark
      clean_lst = []
      names = df.columns
      for i in names:
          for j in df[i]:
              if j == "      ?":
                  clean_lst.append(i)
                  break
```

```
[85]: clean_lst
```

```
[85]: ['pcv', 'wbcc', 'rbcc']
```

```
[86]: # Replace pcv
      df['pcv'] = df['pcv'].replace({"      ?":49})
```

```
[87]: # Replace values
      def re_value(col, new_val):
          df[col] = df[col].replace({"      ?":new_val})
          return
```

```
[88]: col = 'rbcc'
new_val = 8.5
re_value(col, new_val)
```

```
[89]: # Look at the occurrence of the values
df.describe()
```

```
[89]:
```

	age	bp	sg	al	su	rbc	pc	ba	bgr	bu	...	\
count	400	400	400	400	400	400	400	400	400	400	...	
unique	77	11	6	7	7	2	3	3	147	119	...	
top	60	80	1.020	0	0	normal	normal	notpresent	150	?	...	
freq	19	116	106	199	290	353	259	374	44	19	...	

	wbcc	rbcc	htn	dm	cad	appet	pe	ane	class	Unnamed: 24
count	400	400.0	400	400	400	400	400	400	400	1
unique	91	50.0	3	7	5	4	4	3	4	1
top	?	5.0	no	no	no	good	no	no	ckd	notckd
freq	105	130.0	251	257	362	316	322	339	248	1

[4 rows x 25 columns]

```
[118]: # Look for other columns with ? mark
clean_lst = []
names = df.columns
for i in names:
    for j in df[i]:
        if j == "?":
            clean_lst.append(i)
            break
```

```
[122]: # Look for other columns with ? mark
clean_lst2 = []
names = df.columns
for i in names:
    for j in df[i]:
        if j == "      ?":
            clean_lst2.append(i)
            break
```

```
[123]: clean_lst2
```

```
[123]: ['wbcc']
```

```
[119]: clean_lst
```

```
[119]: ['bu']
```

```
[92]: # Replace values
def re_value(col, new_val):
    df[col] = df[col].replace({"?":new_val})
    return
```

```
[93]: #{age:55, bp:80,sg:1.010, al:1,su:0,pc:normal(0)abnormal=1,ba:
    ↪0=present,1=notpresent
#sc:9.0,sod:135,pot:8.0,hemo:13.5,pcv:38,wbc:8000, htn:0=no,1=yes,dm:0=no,1=yes
#cad:0=no,1=yes,appet:poor=1good=0, pe00=no,1=yes, ane00=no,1=yes,rbc:
    ↪normal(0)abnormal=1

new_data = {'age':55, 'bp':80,'sg':1.010, 'al':1,'su':0,
            'sc':9.0,'sod':135,'pot':8.0,'hemo':13.5,'pcv':38,'wbcc':8000}

for k,v in new_data.items():
    col = k
    new_val = v
    new_col = re_value(col, new_val)
```

```
[94]: df['bu'].head(15)
```

```
[94]: 0      36
      1      18
      2      53
      3      56
      4      26
      5      25
      6      54
      7      31
      8      60
      9     107
     10      55
     11      60
     12      72
     13      86
     14      90
Name: bu, dtype: object
```

```
[120]: df['bu'] = df['bu'].replace({"?":90})
```

```
[124]: df['wbcc'] = df['wbcc'].replace({"      ?":8000})
```

```
[95]: # Define cat cols
cat_cols = ['htn', 'dm', 'cad', 'pe', 'ane', 'class', 'rbc', 'pc', 'appet',
    ↪'ba']
#cat_cols1 = ['htn', 'dm', 'cad', 'pe', 'ane']
#cat_cols2 = ['rbc', 'pc']
```

```
#cat_cols3 = ['appet']
#cat_cols4 = ['ba']
```

```
[96]: # Replace cat values with binary values
```

```
for i in cat_cols:
    df[i] = df[i].replace({"?":0})
```

```
[97]: #Replace other extraneous chars
```

```
df['pc'] = df['pc'].replace({"?":0})
df['class'] = df['class'].replace({"ckd":1, "notckd":0, "'ckd\t":1})
df['pe'] = df['pe'].replace({"good":0})
df['appet'] = df['appet'].replace({"no":0, "?":1})
df['ba'] = df['ba'].replace({"?":1})

df['dm'] = df['dm'].replace({" yes":1, "\tno":0, "\tyes":1, "normal":0})
```

```
[98]: df['dm'] = df['dm'].replace({"yes":1, "no":0,})
```

```
[99]: df['class'] = df['class'].replace({"ckd\t":1})
```

```
[115]: # Check for other values in categorical columns
```

```
for i in cat_cols:
    print(f'{i} {df[i].unique()}')
```

```
htn [1 0]
dm [1 0]
cad [0 1]
pe [0 1]
ane [0 1]
class [1 0]
rbc [0 1]
pc [0 1]
appet [0 1]
ba [0 1]
```

```
[101]: df['class'] = df['class'].replace({"no":0})
```

```
[114]: df['htn'] = df['htn'].replace({"no":0, "yes":1})
```

```
[113]: df['rbc'] = df['rbc'].replace({"normal":0, "abnormal":1})
```

```
[103]: df['cad'] = df['cad'].replace({"no":0, "yes":1, "No":0, "\tno":0})
```

```
[104]: df['ane'] = df['ane'].replace({"no":0, "yes":1})
```

```
[105]: df['pe'] = df['pe'].replace({"no":0, "yes":1})
```



```
[106]: df['htn'] = df['htn'].replace({"normal":0, "abnormal":1})
```

```
[107]: df['pc'] = df['pc'].replace({"normal":0, "abnormal":1})
```

```
[108]: df['appet'] = df['appet'].replace({"good":0, "poor":1})
```

```
[109]: df['ba'] = df['ba'].replace({"notpresent":0, "present":1})
```

```
[116]: df.dtypes
```

```
[116]: age      object
      bp      object
      sg      object
      al      object
      su      object
      rbc     int64
      pc     int64
      ba     int64
      bgr     object
      bu     object
      sc     object
      sod     object
      pot     object
      hemo    object
      pcv     object
      wbcc    object
      rbcc    object
      htn     int64
      dm     int64
      cad     int64
      appet  int64
      pe     int64
      ane     int64
      class  int64
      dtype: object
```

```
[111]: df.drop(columns=['Unnamed: 24'], axis=1, inplace=True)
```

```
[125]: # Change df to numeric
      df = df.apply(pd.to_numeric)
```

```
[126]: df.dtypes
```

```
[126]: age      int64
      bp      int64
      sg     float64
      al      int64
```

```

su          int64
rbc         int64
pc          int64
ba          int64
bgr         int64
bu          float64
sc          float64
sod         float64
pot         float64
hemo        float64
pcv         int64
wbcc        int64
rbcc        float64
htn         int64
dm          int64
cad         int64
appet       int64
pe          int64
ane         int64
class       int64
dtype: object

```

```

[127]: # Keep a copy just incase need to revert back to this
df_copy = df
df_copy.to_csv('new_ckd_data2.csv')

```

```

[128]: #Read in the new dataset
df = pd.read_csv('new_ckd_data2.csv', header=0)

```

```

[129]: df.drop("Unnamed: 0", axis=1, inplace=True)

```

```

[40]: # Check values in certain columns

```

2 Perform Data Correlations and Relationship Analysis

```

[130]: # Compute correlation between features
corr_df = df.corr()
corr_df

```

```

[130]:
```

	age	bp	sg	al	su	rbc	pc \
age	1.000000	0.136121	-0.212145	0.111207	0.185521	0.016852	0.108413
bp	0.136121	1.000000	-0.211619	0.147564	0.189561	0.150384	0.156231
sg	-0.212145	-0.211619	1.000000	-0.427115	-0.222632	-0.161723	-0.251734
al	0.111207	0.147564	-0.427115	1.000000	0.262564	0.374484	0.535895
su	0.185521	0.189561	-0.222632	0.262564	1.000000	0.092940	0.190062
rbc	0.016852	0.150384	-0.161723	0.374484	0.092940	1.000000	0.377394

pc	0.108413	0.156231	-0.251734	0.535895	0.190062	0.377394	1.000000
ba	0.043633	0.110164	-0.165956	0.368222	0.119399	0.184402	0.330401
bgr	0.212837	0.149744	-0.311199	0.327390	0.629700	0.154056	0.263195
bu	0.170725	0.182174	-0.320948	0.396791	0.122439	0.237919	0.338142
sc	0.101705	0.139881	-0.283513	0.215704	0.085736	0.143396	0.152913
sod	-0.085662	-0.107184	0.311943	-0.281730	-0.062081	-0.147691	-0.183487
pot	0.051353	0.084987	-0.175673	0.152497	0.195494	0.014188	0.185147
hemo	-0.168518	-0.273970	0.550722	-0.543178	-0.155351	-0.279270	-0.410533
pcv	-0.204758	-0.295150	0.546819	-0.525709	-0.182275	-0.271486	-0.418691
wbcc	0.092800	0.022177	-0.141381	0.191492	0.156649	-0.003471	0.106733
rbcc	-0.190103	-0.215067	0.393612	-0.411869	-0.156812	-0.181685	-0.372117
htn	0.386340	0.266901	-0.438216	0.480675	0.254268	0.140538	0.291719
dm	0.355239	0.226489	-0.458341	0.377038	0.430514	0.145646	0.201032
cad	0.232499	0.084135	-0.186755	0.236254	0.229301	0.111493	0.172295
appet	0.150940	0.175054	-0.259634	0.330264	0.069216	0.160868	0.274985
pe	0.093422	0.056902	-0.289736	0.440612	0.116442	0.199285	0.350227
ane	0.048092	0.195134	-0.265119	0.281546	0.042464	0.107625	0.260566
class	0.222755	0.293693	-0.720910	0.598389	0.294555	0.282642	0.375154

	ba	bgr	bu	...	pcv	wbcc	rbcc	\
age	0.043633	0.212837	0.170725	...	-0.204758	0.092800	-0.190103	
bp	0.110164	0.149744	0.182174	...	-0.295150	0.022177	-0.215067	
sg	-0.165956	-0.311199	-0.320948	...	0.546819	-0.141381	0.393612	
al	0.368222	0.327390	0.396791	...	-0.525709	0.191492	-0.411869	
su	0.119399	0.629700	0.122439	...	-0.182275	0.156649	-0.156812	
rbc	0.184402	0.154056	0.237919	...	-0.271486	-0.003471	-0.181685	
pc	0.330401	0.263195	0.338142	...	-0.418691	0.106733	-0.372117	
ba	1.000000	0.086393	0.149388	...	-0.190211	0.104748	-0.188575	
bgr	0.086393	1.000000	0.131373	...	-0.269870	0.114906	-0.195127	
bu	0.149388	0.131373	1.000000	...	-0.519610	0.034946	-0.439877	
sc	0.039119	0.087844	0.589841	...	-0.333726	-0.018594	-0.285009	
sod	-0.078859	-0.172496	-0.295783	...	0.350864	0.012154	0.277338	
pot	-0.012290	0.119964	0.276549	...	-0.147251	-0.084916	-0.089104	
hemo	-0.202931	-0.263433	-0.523121	...	0.849999	-0.139315	0.610464	
pcv	-0.190211	-0.269870	-0.519610	...	1.000000	-0.171419	0.643391	
wbcc	0.104748	0.114906	0.034946	...	-0.171419	1.000000	-0.165621	
rbcc	-0.188575	-0.195127	-0.439877	...	0.643391	-0.165621	1.000000	
htn	0.089046	0.369813	0.367661	...	-0.567821	0.116827	-0.512260	
dm	0.080070	0.500175	0.300813	...	-0.451991	0.144101	-0.378652	
cad	0.162395	0.212723	0.217758	...	-0.292816	0.006923	-0.288533	
appet	0.149126	0.175960	0.266698	...	-0.398166	0.142714	-0.374219	
pe	0.134732	0.101941	0.332424	...	-0.391370	0.123381	-0.307671	
ane	0.052208	0.127844	0.427274	...	-0.513960	0.034059	-0.377516	
class	0.186871	0.402785	0.372208	...	-0.698160	0.177571	-0.488995	

	htn	dm	cad	appet	pe	ane	class
age	0.386340	0.355239	0.232499	0.150940	0.093422	0.048092	0.222755

bp	0.266901	0.226489	0.084135	0.175054	0.056902	0.195134	0.293693
sg	-0.438216	-0.458341	-0.186755	-0.259634	-0.289736	-0.265119	-0.720910
al	0.480675	0.377038	0.236254	0.330264	0.440612	0.281546	0.598389
su	0.254268	0.430514	0.229301	0.069216	0.116442	0.042464	0.294555
rbc	0.140538	0.145646	0.111493	0.160868	0.199285	0.107625	0.282642
pc	0.291719	0.201032	0.172295	0.274985	0.350227	0.260566	0.375154
ba	0.089046	0.080070	0.162395	0.149126	0.134732	0.052208	0.186871
bgr	0.369813	0.500175	0.212723	0.175960	0.101941	0.127844	0.402785
bu	0.367661	0.300813	0.217758	0.266698	0.332424	0.427274	0.372208
sc	0.240687	0.185612	0.174665	0.154860	0.171187	0.226633	0.294956
sod	-0.306791	-0.277076	-0.210044	-0.167510	-0.148358	-0.199578	-0.379096
pot	0.059907	0.105091	-0.018690	0.017377	0.049738	0.096755	0.224031
hemo	-0.580489	-0.462453	-0.289364	-0.376544	-0.370653	-0.558788	-0.699533
pcv	-0.567821	-0.451991	-0.292816	-0.398166	-0.391370	-0.513960	-0.698160
wbcc	0.116827	0.144101	0.006923	0.142714	0.123381	0.034059	0.177571
rbcc	-0.512260	-0.378652	-0.288533	-0.374219	-0.307671	-0.377516	-0.488995
htn	1.000000	0.608118	0.325479	0.345070	0.371026	0.347802	0.590438
dm	0.608118	1.000000	0.271172	0.325134	0.308463	0.183686	0.559060
cad	0.325479	0.271172	1.000000	0.156104	0.172295	0.047700	0.236088
appet	0.345070	0.325134	0.156104	1.000000	0.417055	0.254942	0.393341
pe	0.371026	0.308463	0.172295	0.417055	1.000000	0.207025	0.375154
ane	0.347802	0.183686	0.047700	0.254942	0.207025	1.000000	0.325396
class	0.590438	0.559060	0.236088	0.393341	0.375154	0.325396	1.000000

[24 rows x 24 columns]

```
[131]: # Define seaborn parameters

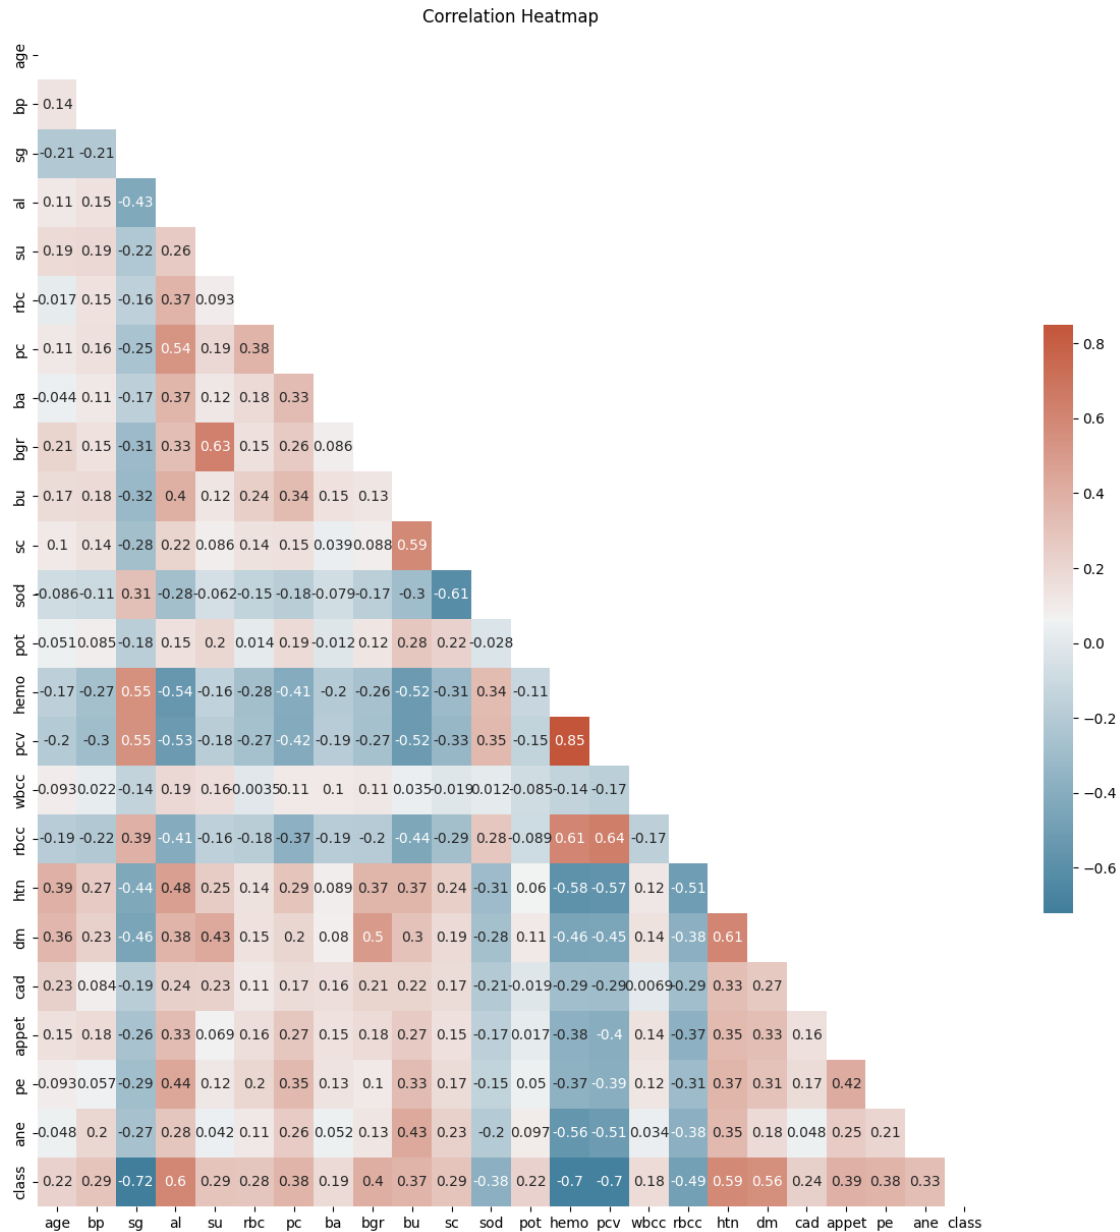
#Generate a cubehelix heatmap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

#Create a mask to display the upper triangle
mask = np.triu(np.ones_like(corr_df, dtype=bool))

# Set up the figure bounds

f, ax = plt.subplots(figsize=(15,15))

sns.heatmap(corr_df, mask=mask, cmap=cmap, cbar_kws={"shrink": .5},
            annot=True)
plt.title('Correlation Heatmap')
plt.savefig("sns-heatmap.png")
plt.show()
```



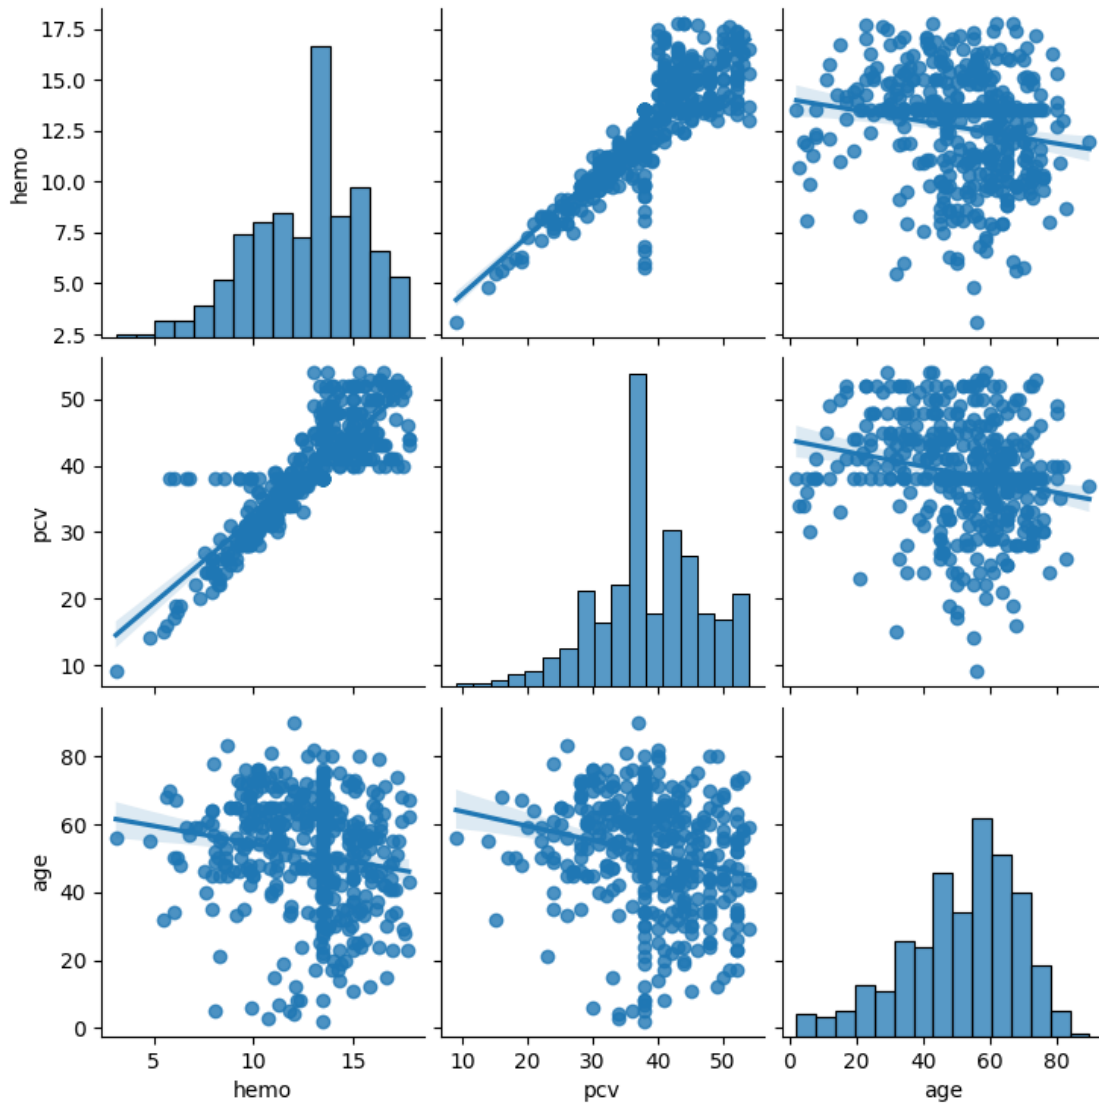
```
[43]: # There doesn't seem to be many positive correlation between features
      # Person with anemia is expected to have a lower hemoglobin and low red blood
      # cell counts
      # Check the analysis
```

```
[44]: #Look at some variables with seaborn pairplots
```

```
[134]: plt.figure(figsize=(14,10))
```

```
sns.pairplot(df[['hemo', 'pcv', 'age']], kind="reg")
plt.show()
```

<Figure size 1400x1000 with 0 Axes>

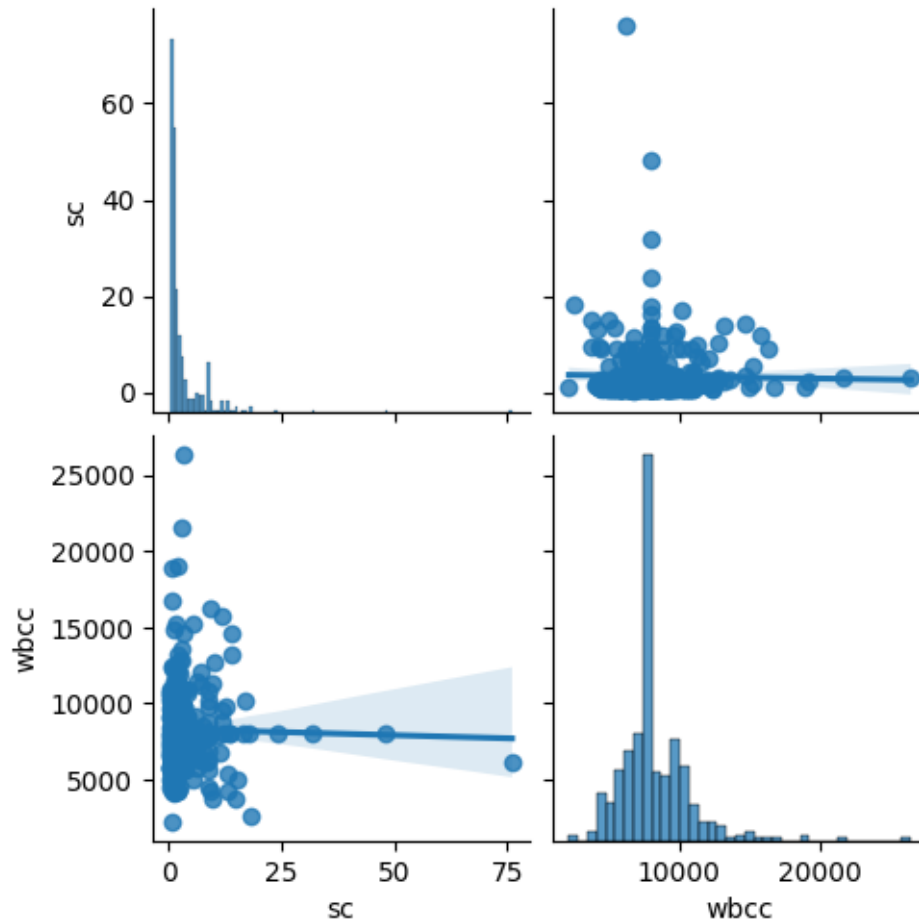


```
[135]: # Look at serum creatinine and potassium and sodium levels
# There appears to be a linear relationship between sc and hemoglobin
```

```
plt.figure(figsize=(10,10))

sns.pairplot(df[['sc', 'wbcc']], kind="reg")
plt.show()
```

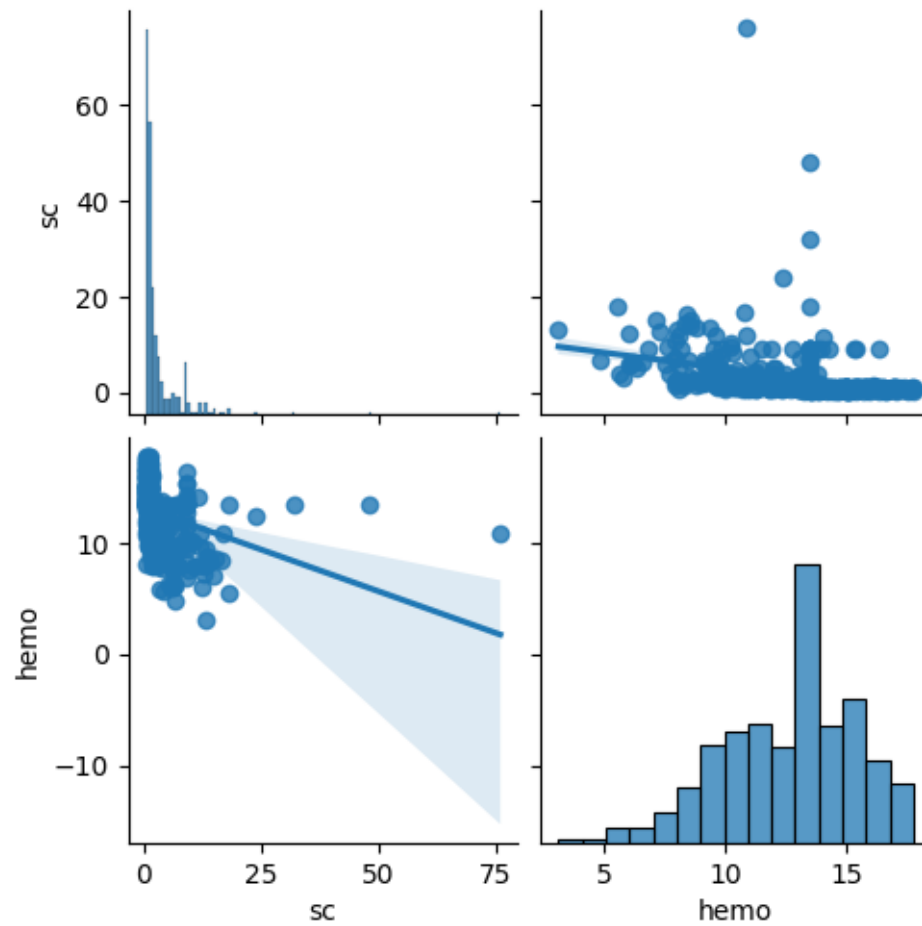
<Figure size 1000x1000 with 0 Axes>



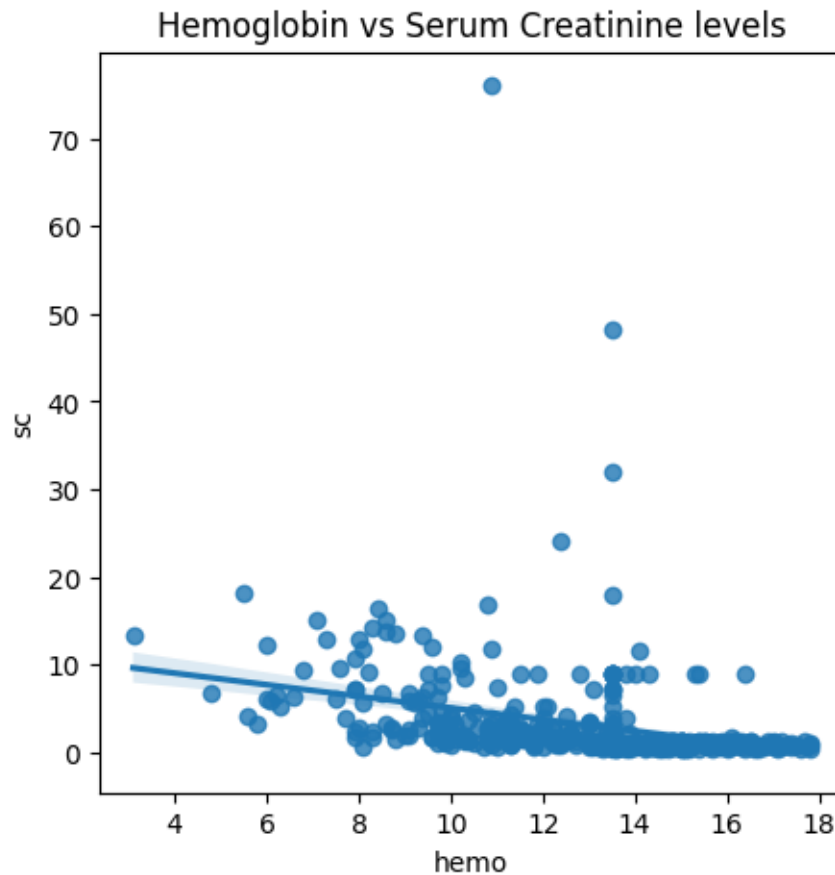
```
[136]: # Look at serum creatinine and hemoglobin  
# There appears to be a linear relationship between sc and hemoglobin
```

```
plt.figure(figsize=(14,10))  
sns.pairplot(df[['sc', 'hemo']], kind="reg")  
plt.show()
```

<Figure size 1400x1000 with 0 Axes>

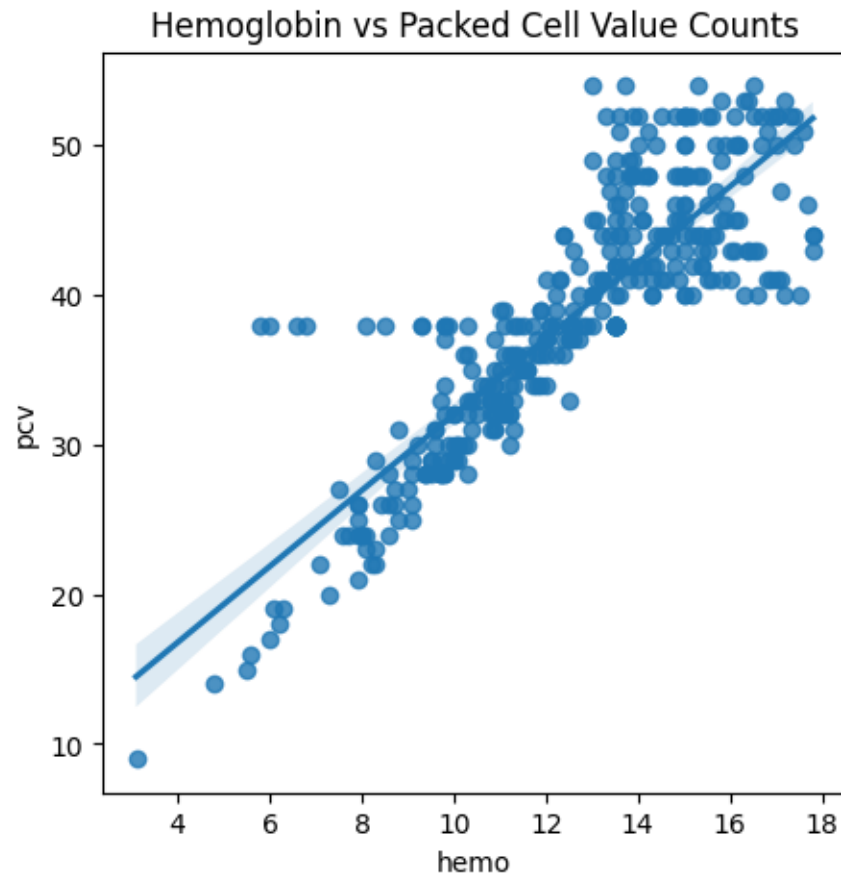


```
[137]: # look at a specific pair
plt.figure(figsize=(5,5))
xlabel = "Hemoglobin Levels"
ylabel = "Serum Creatinine Levels"
sns.regplot(x=df["hemo"], y=df["sc"])
plt.title('Hemoglobin vs Serum Creatinine levels')
plt.show()
```





```
[138]: # look at a specific pair
# This is expected as hemoglobin counts rise so do the number of red blood cells
# This is a natural occurrence and in itself does not show propensity towards CKD
# Pvc can be deleted from the feature set for the model.

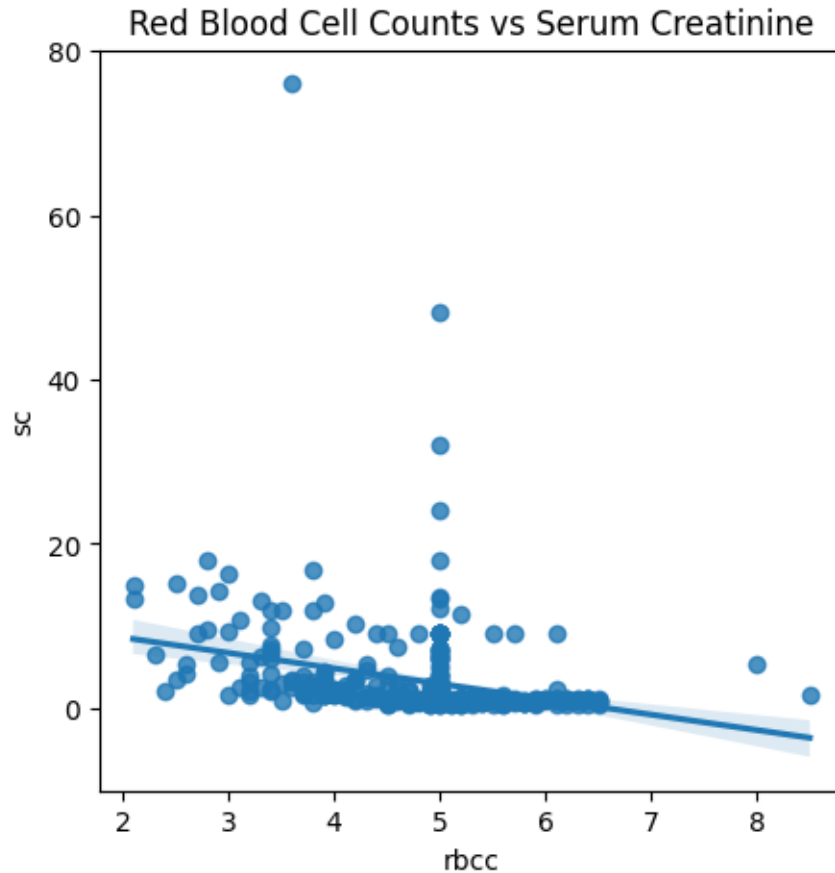
plt.figure(figsize=(5,5))
xlabel = "Hemoglobin Levels"
ylabel = "Packed Cell Value Counts"
sns.regplot(x=df["hemo"], y=df["pcv"])
plt.title('Hemoglobin vs Packed Cell Value Counts')
plt.show()
```



```
[139]: #Does age factor into the equation?  
  
plt.figure(figsize=(5,5))  
xlabel = "Hemoglobin Levels"  
ylabel = "Age"  
sns.regplot(x=df["hemo"], y=df["age"])  
plt.title('Hemoglobin vs Age')  
plt.show()
```



```
[140]: # Appears to have a linear relationship
# As rbcc counts increase, serum creatinine decrease
# makes sense because higher creatinine levels are associated with
# lower hemoglobin counts which indicates anemia
plt.figure(figsize=(5,5))
xlabel = "Red Blood Cell Counts"
ylabel = "Serum Creatinine"
sns.regplot(x=df["rbcc"], y=df["sc"])
plt.title('Red Blood Cell Counts vs Serum Creatinine')
plt.show()
```

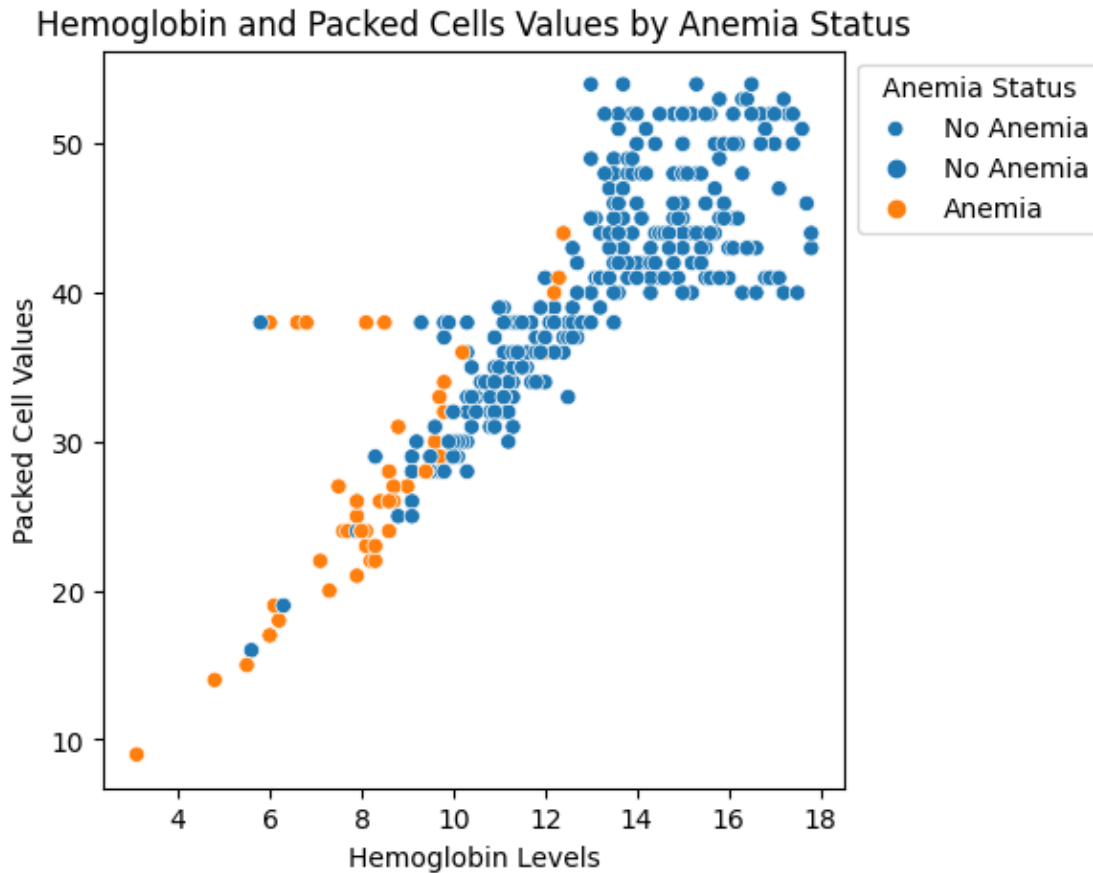


```
[52]: # Analysis the data using Scatterplots to compare features
```

```
[142]: #Create a scatter plot function using seaborn
def sns_scatter_plt(X,Y,H, xlab, ylab, Title):
    plt.figure(figsize=(5,5))
    ax = sns.scatterplot(x=df[X],y=df[Y],hue=df[H])
    plt.xlabel(xlab)
    plt.ylabel(ylab)
    plt.title(Title)
    plt.legend()
    return (ax)
```

```
[145]: #Relationship between hemoglobin counts and Packed Cell Volumes by Anemia Status
# Expect the lower the hemoglobin and packed cell values to indicate anemia
X='hemo'
Y='pcv'
H='ane'
xlab='Hemoglobin Levels'
ylab='Packed Cell Values'
```

```
Title = 'Hemoglobin and Packed Cells Values by Anemia Status'
ax = sns_scatter_plt(X,Y,H, xlab, ylab, Title)
ax.legend(title='Anemia Status', bbox_to_anchor=(1, 1), labels=['No Anemia', 'No Anemia', 'Anemia'])
plt.show()
```

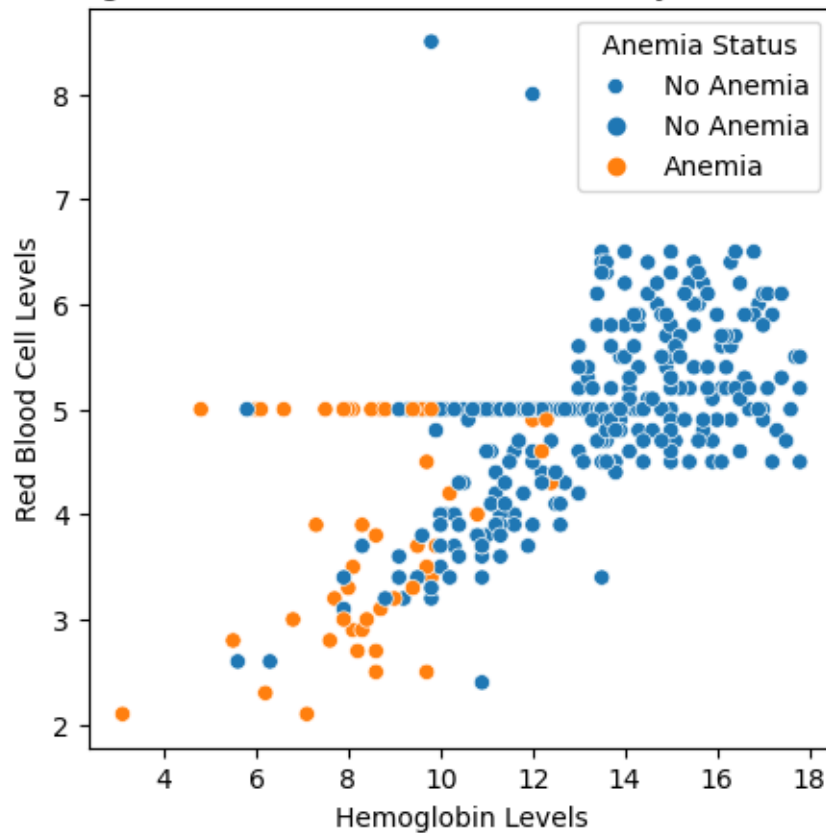


```
[146]: #Relationship between hemoglobin counts and Red Blood cell Volumes by Anemia Status
# Expect the lower the hemoglobin and Red Blood cell values to indicate anemia

X='hemo'
Y='rbcc'
H='ane'
xlab='Hemoglobin Levels'
ylab='Red Blood Cell Levels'
Title = 'Hemoglobin and Red Blood Cell Values by Anemia Status'
ax = sns_scatter_plt(X,Y,H, xlab, ylab, Title)
ax.legend(title='Anemia Status', bbox_to_anchor=(1, 1), labels=['No Anemia', 'No Anemia', 'Anemia'])
```

```
plt.show()
```

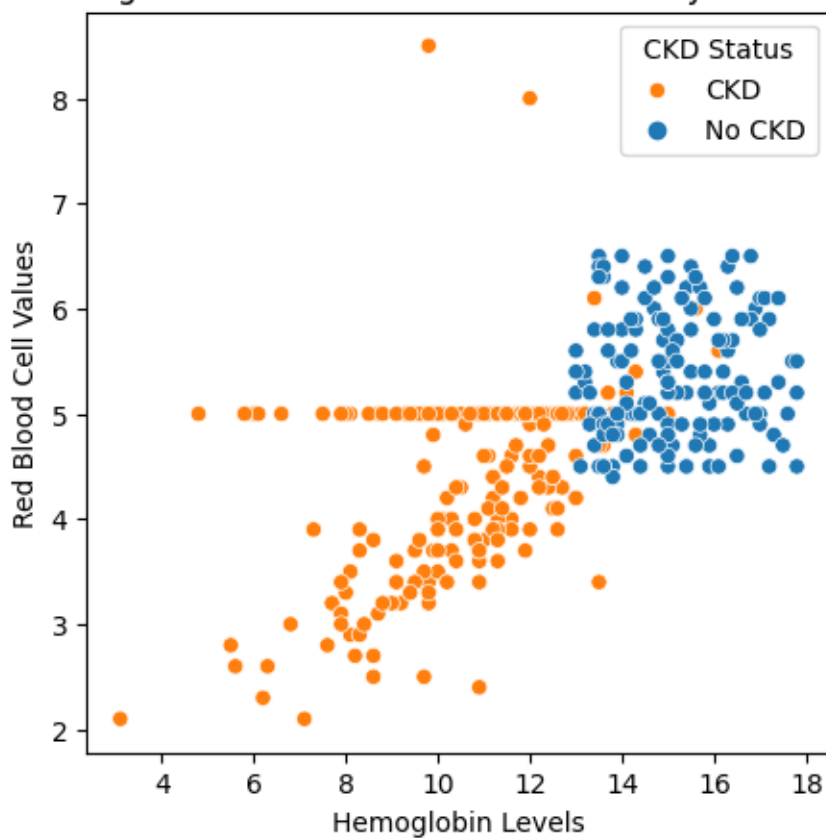
Hemoglobin and Red Blood Cell Values by Anemia Status



```
[147]: # Relationship between hemoglobin counts and Red Blood Cell Values by CKD Status
# With CKD comes anemia, expect lower hemoglobin and red blood cell values.
X='hemo'
Y='rbcc'
H='class'
xlab='Hemoglobin Levels'
ylab='Red Blood Cell Values'
Title = 'Hemoglobin and Red Blood Cells Values by CKD Status'

ax = sns_scatterplt(X,Y,H, xlab, ylab, Title)
ax.legend(title='CKD Status', bbox_to_anchor=(1, 1), labels=['CKD', 'No CKD'])
plt.show()
```

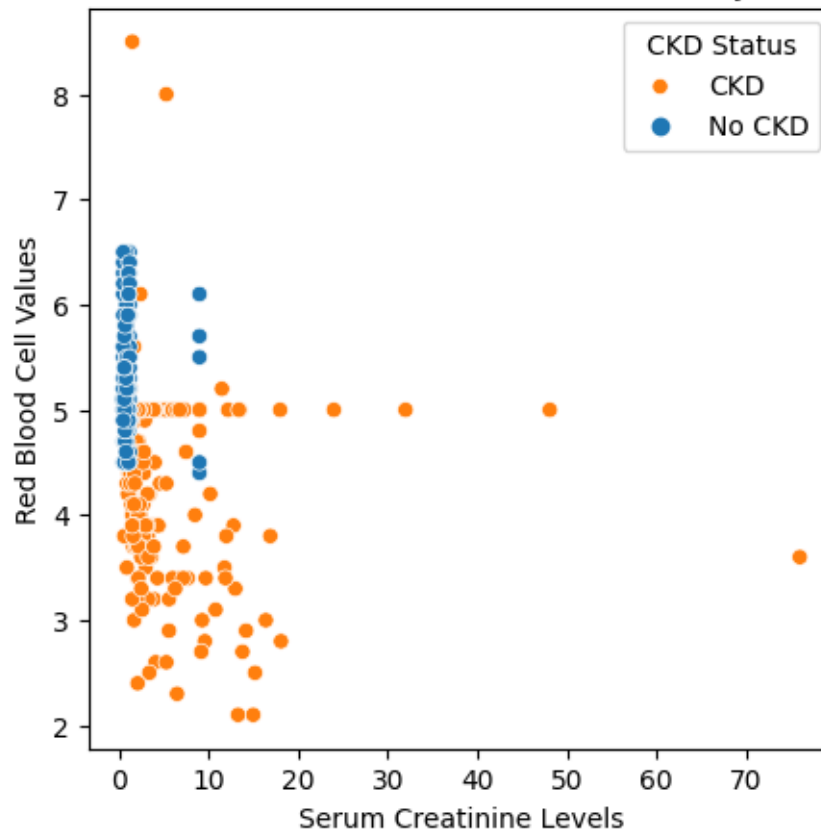
Hemoglobin and Red Blood Cells Values by CKD Status



```
[281]: # Relationship between Serum Creatinine and Red Blood Cell Values by CKD Status
# With CKD comes anemia, expect lower hemoglobin and red blood cell values.
X='sc'
Y='rbcc'
H='class'
xlab='Serum Creatinine Levels'
ylab='Red Blood Cell Values'
Title = 'Serum Creatinine and Red Blood Cells Values by CKD Status'

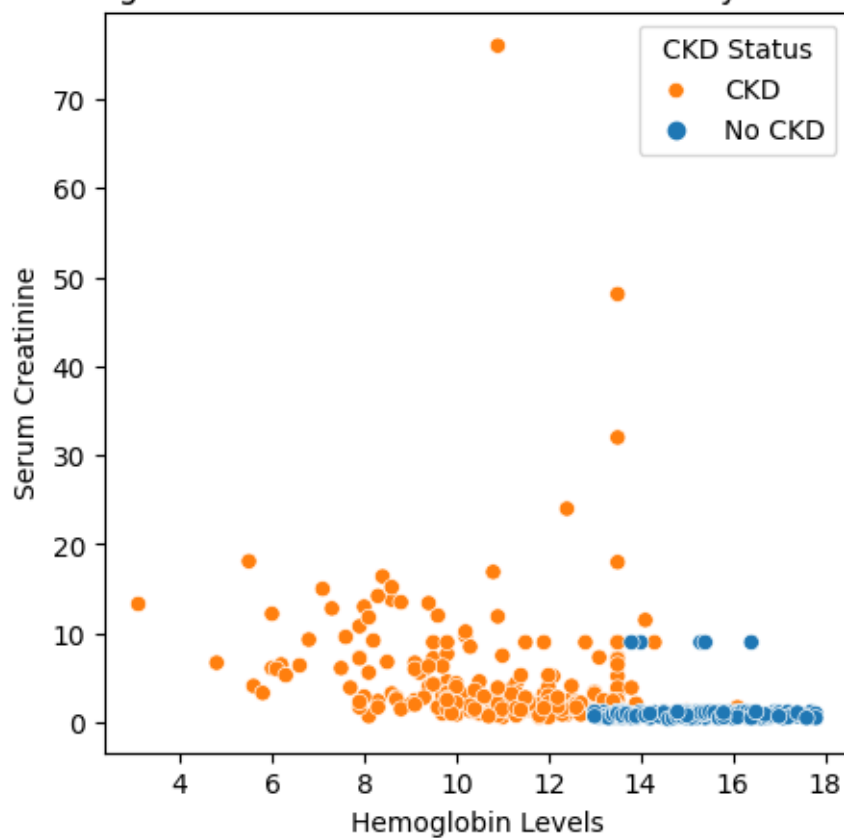
ax = sns_scatter(X,Y,H, xlab, ylab, Title)
ax.legend(title='CKD Status', bbox_to_anchor=(1, 1), labels=['CKD', 'No CKD'])
plt.show()
```

Serum Creatinine and Red Blood Cells Values by CKD Status




```
[148]: # Relationship between Serum Creatinine and Red Blood Cell Values by CKD Status
# With CKD comes anemia, expect lower hemoglobin and red blood cell values.
X='hemo'
Y='sc'
H='class'
xlab='Hemoglobin Levels'
ylab='Serum Creatinine'
Title = 'Hemoglobin and Serum Creatinine Values by CKD Status'
ax = sns_scatter_plot(X,Y,H, xlab, ylab, Title)
ax.legend(title='CKD Status', bbox_to_anchor=(1, 1), labels=['CKD', 'No CKD'])
plt.show()
```


Hemoglobin and Serum Creatinine Values by CKD Status



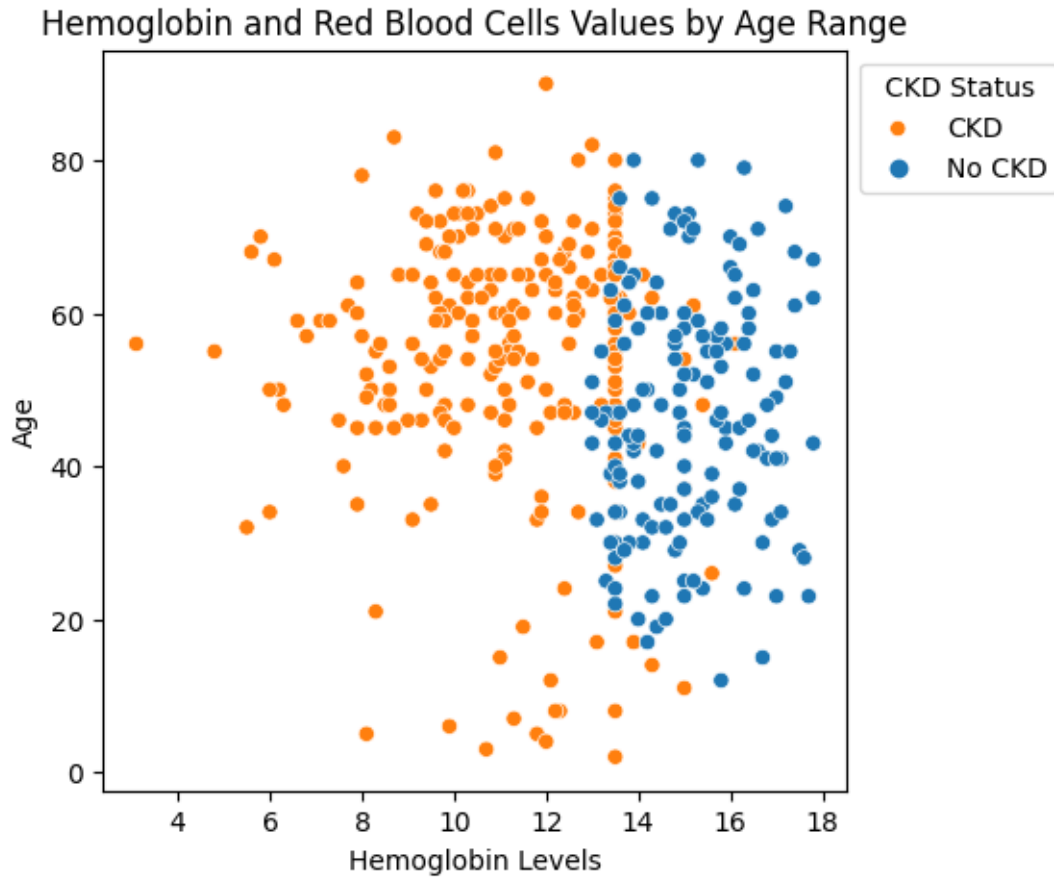
```
[149]: # Relationship between Serum Creatinine and Red Blood Cell Values by CKD Status
# With CKD comes anemia, expect lower hemoglobin and red blood cell values.
X='hemo'
Y='sc'
H='ane'
xlab='Hemoglobin Levels'
ylab='Serum Creatinine'
Title = 'Hemoglobin and Red Blood Cells Values by Anemia Status'
ax = sns_scatter(X,Y,H, xlab, ylab, Title)
ax.legend(title='Anemia Status', bbox_to_anchor=(1, 1), labels=['No Anemia', 'Anemia'])
plt.show()
```

Hemoglobin and Red Blood Cells Values by Anemia Status



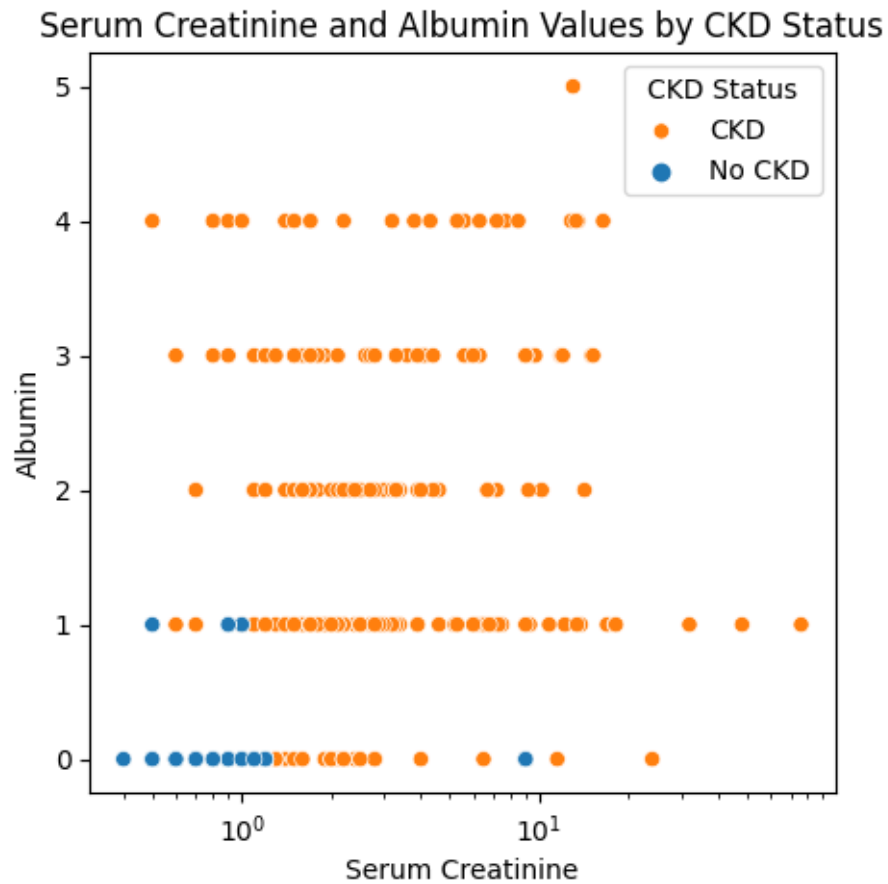
```
[150]: # Check the relationships by age
# Relationship between age and anemia (low hemoglobin values and low rbcc)

X='hemo'
Y='age'
H='class'
xlab='Hemoglobin Levels'
ylab='Age'
Title = 'Hemoglobin and Red Blood Cells Values by Age Range'
ax = sns_scatter(X,Y,H, xlab, ylab, Title)
ax.legend(title='CKD Status', bbox_to_anchor=(1, 1), labels=['CKD', 'No CKD'])
plt.show()
```

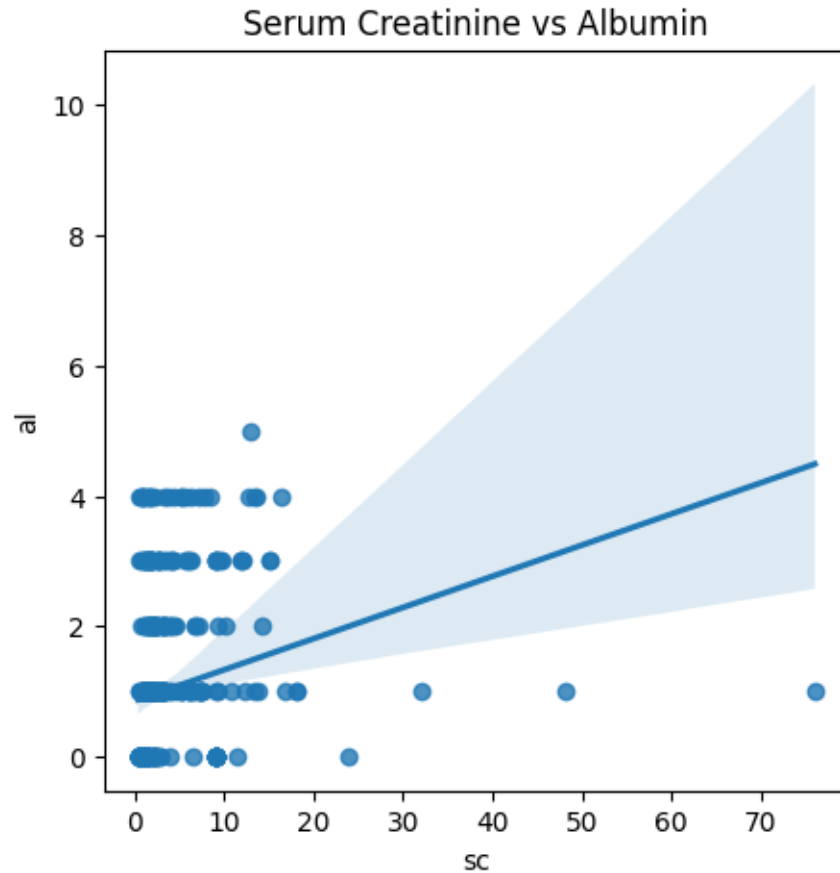


```
[158]: # Check the relationships by albumin
# Relationship between albumin and serum creatinine

X='sc'
Y='al'
H='class'
xlab='Serum Creatinine'
ylab='Albumin'
Title = 'Serum Creatinine and Albumin Values by CKD Status'
ax = sns_scatter(X,Y,H, xlab, ylab, Title)
ax.legend(title='CKD Status', bbox_to_anchor=(1, 1), labels=['CKD', 'No CKD'])
ax.set_xscale("log")
plt.show()
```



```
[161]: plt.figure(figsize=(5,5))
xlabel = "Serum Creatinine"
ylabel = "Albumin"
sns.regplot(x=df["sc"], y=df["al"])
plt.title('Serum Creatinine vs Albumin')
plt.show()
```



```
[162]: fig = plt.figure(figsize=(10,4))
title = fig.suptitle('Hypertension and Diabetes Counts By CKD Status',
    ↳ fontsize=12, fontweight="bold")

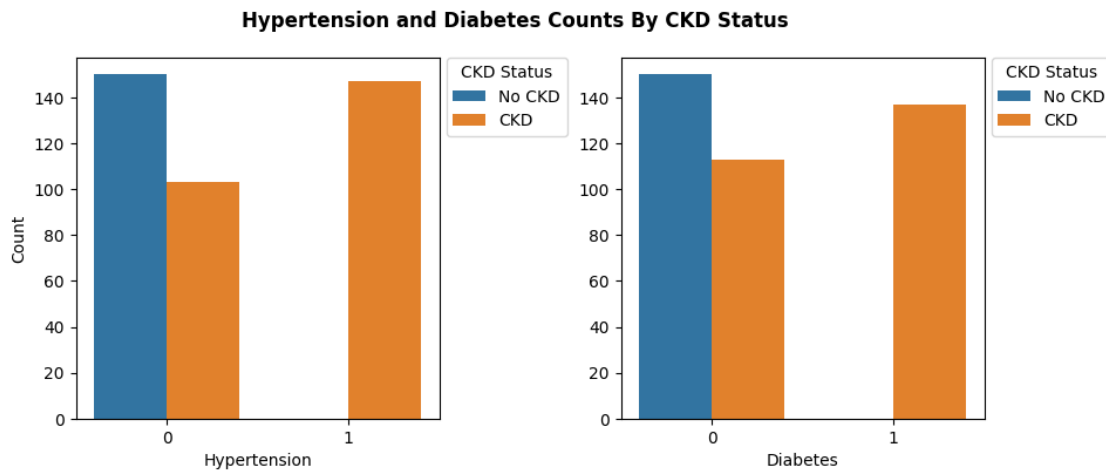
fig.subplots_adjust(top=0.88, wspace=0.5)

# Countplot 1
ax1 = ax1 = fig.add_subplot(1,2,1)

ax1 = sns.countplot(x='htn',hue='class',data=df)
ax1.set_xlabel('Hypertension')
ax1.set_ylabel("Count")
plt.legend(title='CKD Status', labels=['No CKD', 'CKD'],bbox_to_anchor=(1.02,
    ↳ 1), loc='upper left', borderaxespad=0)

# Countplot 2
ax2 = ax2 = fig.add_subplot(1,2,2)
```

```
ax2 = sns.countplot(x='dm',hue='class',data=df)
ax2.set_xlabel('Diabetes')
ax2.set_ylabel("")
plt.legend(title='CKD Status', labels=['No CKD', 'CKD'],bbox_to_anchor=(1.02,1), loc='upper left', borderaxespad=0)
plt.show()
```



3 Feature Selection

```
[167]: #Create features and target sets
# Dependent variable is the CKD status or 'class'
# y is the class column, last column in the dataset
```

```
[168]: X = df.iloc[:, :-1] #features
y = df.iloc[:, -1:] #target
```

```
[169]: # Normalize the dataset
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_scale = scaler.fit_transform(X)
y_scale = scaler.fit_transform(y)
```

```
[170]: # Create the training and test datasets
# Train Test Split
from sklearn.model_selection import train_test_split
```

```
[171]: X_train, X_test, y_train, y_test = train_test_split(X_scale, y_scale, test_size=0.2, random_state = 0)
```

```
[172]: # Only 400 obs in df, check the shape of X_train and X_test
X_train.shape, X_test.shape
```

```
[172]: ((320, 23), (80, 23))
```

```
[173]: #Conduct Feature Engineering
```

```
[174]: #Use RFE from SKlearn
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
```

```
[175]: # Build model
rf = RandomForestRegressor()
RFE(estimator=LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1,
    ↪normalize=False),
    n_features_to_select=None, step=1, verbose=0)

#Fit the model
rf.fit(X_train,y_train.ravel())

#Check the score using R-squared
rf.score(X_test, y_test.ravel())
```

```
[175]: 0.9285989010989011
```

```
[176]: # View the computed coefficients by viewing the .feature_importances object
#
rf_df = pd.DataFrame(
    zip(X.columns, abs(rf.feature_importances_)),
    columns=["feature", "weight"],
    ).sort_values("weight").reset_index(drop=True)

rf_df
```

```
[176]:   feature  weight
0      ane  0.000000
1      cad  0.000000
2       ba  0.000000
3       pc  0.000118
4       su  0.000519
5     wbcc  0.001393
6      rbc  0.001446
7      age  0.001632
8       pe  0.002120
9       bp  0.002197
10    appet  0.004363
```

11	bgr	0.008341
12	rbcc	0.013321
13	bu	0.014456
14	dm	0.014624
15	pot	0.014970
16	htn	0.015105
17	sod	0.016455
18	sc	0.017323
19	al	0.050762
20	sg	0.173735
21	hemo	0.254775
22	pcv	0.392347

```
[177]: # Feature Selection using RFECV
```

```
[178]: from sklearn.feature_selection import RFECV
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import StratifiedKFold
```

```
[179]: #Build model and fit the model
```

```
[180]: min_features_to_select = 3

clf = LogisticRegression()
cv = StratifiedKFold(5)

rfecv = RFECV(
    estimator=clf,
    step=1,
    cv=cv,
    scoring="accuracy",
    min_features_to_select=min_features_to_select,
    n_jobs=2,
)
rfecv.fit(X_train, y_train.ravel())

print(f"Optimal number of features: {rfecv.n_features_}")
```

Optimal number of features: 13

```
[181]: #Plot the features vs cross-validation scores
```

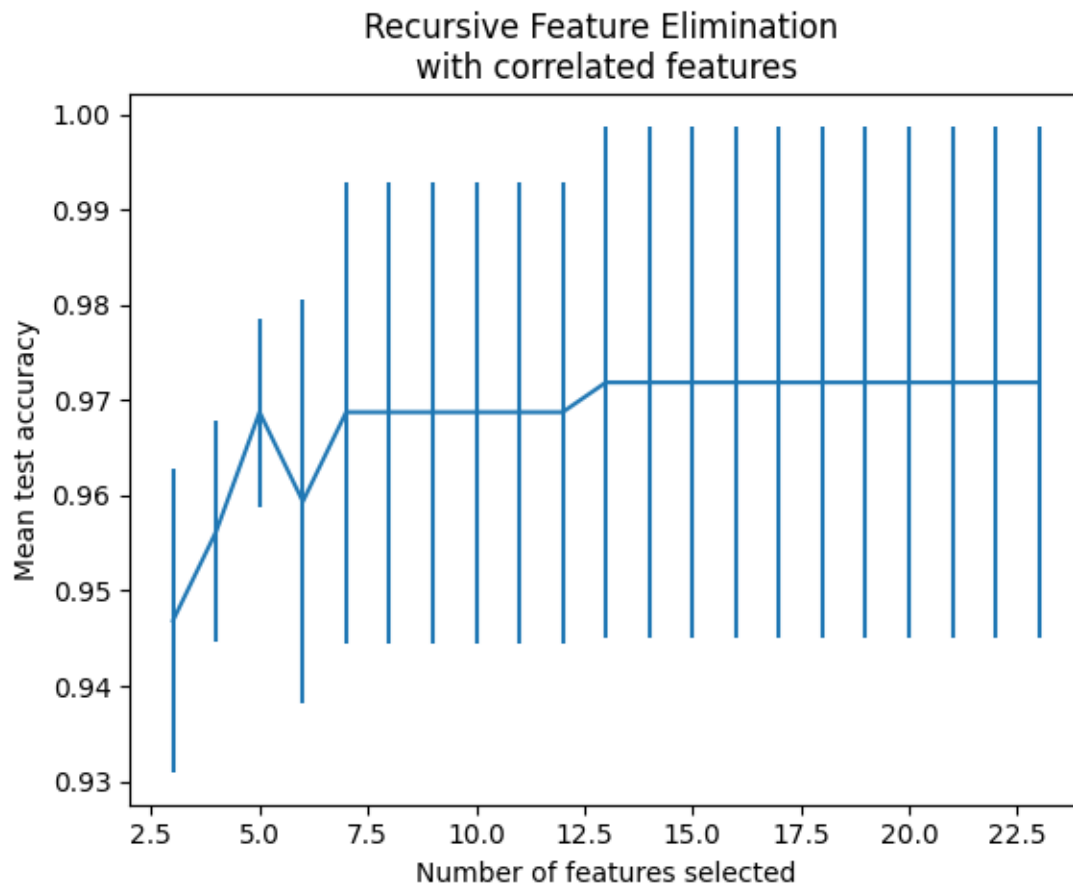
```
n_scores = len(rfecv.cv_results_["mean_test_score"])
plt.figure()
plt.xlabel("Number of features selected")
plt.ylabel("Mean test accuracy")
plt.errorbar(
```



```

range(min_features_to_select, n_scores + min_features_to_select),
rfecv.cv_results_["mean_test_score"],
yerr=rfecv.cv_results_["std_test_score"],
)
plt.title("Recursive Feature Elimination \nwith correlated features")
plt.show()

```



```

[182]: # View the selected features along with the weights from the RFE model
#
rfecv_df = pd.DataFrame(
    zip(X.columns, rfecv.support_, rf_df['weight']),
    columns=["feature", "mask", 'rf weights'],
    ).sort_values("rf weights").reset_index(drop=True)

rfecv_df

```

```

[182]:   feature  mask  rf weights
0     age  False  0.000000
1      bp   True  0.000000

```

2	sg	True	0.000000
3	al	True	0.000118
4	su	False	0.000519
5	rbc	True	0.001393
6	pc	True	0.001446
7	ba	False	0.001632
8	bgr	True	0.002120
9	bu	False	0.002197
10	sc	False	0.004363
11	sod	False	0.008341
12	pot	False	0.013321
13	hemo	True	0.014456
14	pcv	True	0.014624
15	wbcc	False	0.014970
16	rbcc	True	0.015105
17	htn	True	0.016455
18	dm	True	0.017323
19	cad	False	0.050762
20	appet	True	0.173735
21	pe	True	0.254775
22	ane	False	0.392347

```
[183]: # Create the feature dataset from the feature estimator model
# to input to the prediction models
```

```
[184]: #Copy the old df with all the columns from both models
col = rfecv_df['feature'].to_list()
```

```
[185]: new_df = df[col]
```

```
[186]: #add the class column
new_df['class'] = df['class']
```

```
[187]: # select only the feature columns selected by RFECV
rfecv_cols = rfecv_df[rfecv_df['mask'] == True]
```

```
[188]: rfecv_cols
```


```
[188]:
```

	feature	mask	rf weights
1	bp	True	0.000000
2	sg	True	0.000000
3	al	True	0.000118
5	rbc	True	0.001393
6	pc	True	0.001446
8	bgr	True	0.002120
13	hemo	True	0.014456
14	pcv	True	0.014624

16	rbcc	True	0.015105
17	htn	True	0.016455
18	dm	True	0.017323
20	appet	True	0.173735
21	pe	True	0.254775

```
[189]: #Plot the selected features by weight
#Use a barplot
```

```
sns.barplot(x=rfecv_cols['feature'], y = rfecv_cols['rf weights'])
plt.title('Selected Feature from RFECV with Importance Weights from RFE')
plt.xlabel('Features')
plt.ylabel('Feature Weights from RFE')
plt.show()
```



```
[ ]:
```

```
[190]: #add class to the list of best features
best_feature_lst = rfecv_cols['feature'].to_list()
best_feature_lst.append('class')
```

```
[191]: feature_df = new_df[best_feature_lst]
```

```
[192]: # df for prediction models
feature_df
```

```
[192]:
```

	bp	sg	al	rbc	pc	bgr	hemo	pcv	rbcc	htn	dm	appet	pe	class
0	80	1.020	1	0	0	121	15.4	44	5.2	1	1	0	0	1
1	50	1.020	4	0	0	150	11.3	38	5.0	0	0	0	0	1
2	80	1.010	2	0	0	423	9.6	31	5.0	0	1	1	0	1
3	70	1.005	4	0	1	117	11.2	32	3.9	1	0	1	1	1
4	80	1.010	2	0	0	106	11.6	35	4.6	0	0	0	0	1
...
395	80	1.020	0	0	0	140	15.7	47	4.9	0	0	0	0	0
396	70	1.025	0	0	0	75	16.5	54	6.2	0	0	0	0	0
397	80	1.020	0	0	0	100	15.8	49	5.4	0	0	0	0	0
398	60	1.025	0	0	0	114	14.2	51	5.9	0	0	0	0	0
399	80	1.025	0	0	0	131	15.8	53	6.1	0	0	0	0	0

[400 rows x 14 columns]

```
[ ]: #Add the serum-column column for an additional test
```

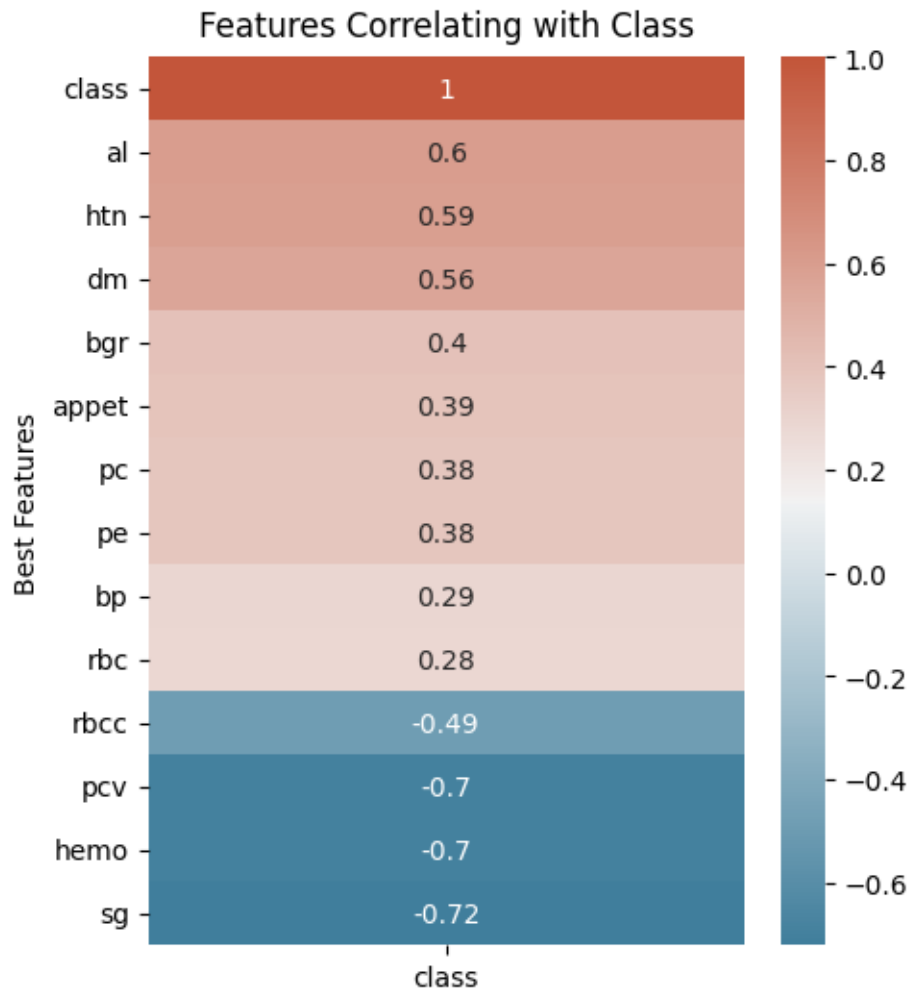
```
[296]: new_test_data = new_df[best_feature_lst].copy()
```

```
[297]: new_test_data['sc'] = new_df['sc'].copy()
```

```
[193]: # Plot the correlation with CKD Status aka Class variable
# Compute the correlation matrix
corr = feature_df.corr()

plt.figure(figsize=(5, 6))

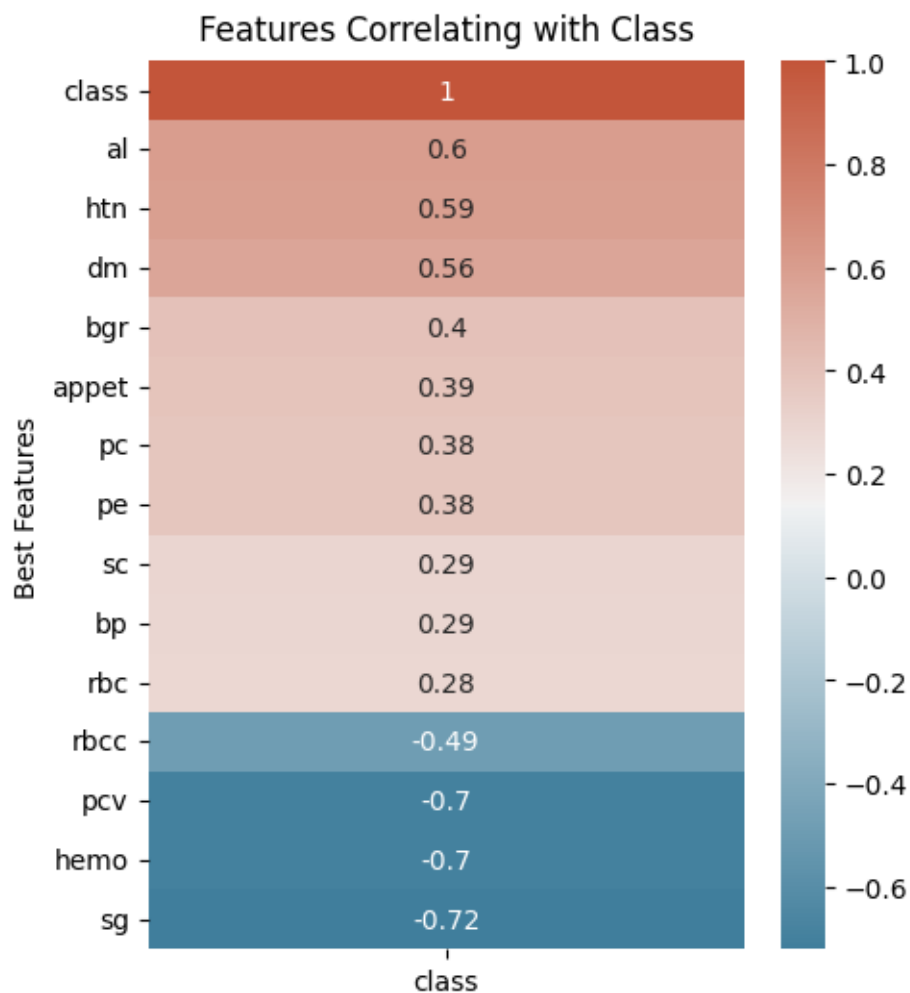
ax = sns.heatmap(feature_df.corr()[['class']].sort_values(by='class',
↪ascending=False), annot=True, cmap=cmap)
ax.set_title('Features Correlating with Class', fontdict={'fontsize':12}, pad=8)
plt.xticks(rotation=0)
plt.ylabel('Best Features')
plt.show()
```



```
[300]: # Plot the correlation with CKD Status using the new_test_set with the sc values
# Compute the correlation matrix
corr = new_test_data.corr()

plt.figure(figsize=(5, 6))

ax = sns.heatmap(new_test_data.corr()[['class']].sort_values(by='class',
    ↪ascending=False), annot=True, cmap=cmap)
ax.set_title('Features Correlating with Class', fontdict={'fontsize':12}, pad=8)
plt.yticks(rotation=0)
plt.ylabel('Best Features')
plt.show()
```



4 Prepare the Training, Test, and Validation Sets

```
[194]: # Read in dataset as validation set
val_df = pd.read_csv('val_set.csv', usecols= best_feature_lst, header=0)

[195]: val_df.head()
```

```
[195]:
```

	bp	sg	al	rbc	pc	bgr	hemo	pcv	rbcc	htn	dm	appet	pe	\
0	80.0	1.025	0.0	0.0	0.0	85.0	15.6	44	6.3	0.0	0.0	1.0	0.0	
1	70.0	1.015	3.0	0.0	0.0	253.0	10.9	31	3.4	1.0	1.0	0.0	1.0	
2	80.0	1.025	0.0	0.0	0.0	119.0	13.9	49	5.0	0.0	0.0	1.0	0.0	
3	80.0	1.020	0.0	0.0	0.0	99.0	13.6	44	6.4	0.0	0.0	1.0	0.0	
4	80.0	1.025	0.0	0.0	0.0	118.0	14.8	45	4.7	0.0	0.0	1.0	0.0	

class

```

0    0.0
1    1.0
2    0.0
3    0.0
4    0.0

```

```

[196]: # sample the validation set
val_set = val_df.sample(n = 60)

```

```

[303]: X2 = feature_df.iloc[:, :-1] #features
y2 = feature_df.iloc[:, -1:] #target
val_set_test = val_set.iloc[:, :-1] #validation set
val_set_class = val_set.iloc[:, -1:] #validation class column

```

```

[198]: #View the value counts for the test set
y2.value_counts()

```

```

[198]: class
1      250
0      150
dtype: int64

```

```

[304]: # Normalize the datasets

scaler = MinMaxScaler()
X_scale2 = scaler.fit_transform(X2)
y_scale2 = scaler.fit_transform(y2)
val_scale_test = scaler.fit_transform(val_set_test)
val_scale_class = scaler.fit_transform(val_set_class)

```

```

[200]: #Create training and test sets
X_train2, X_test2, y_train2, y_test2 = train_test_split(X_scale2, y_scale2,
↳ test_size = 0.2, random_state = 0)

```

```

[201]: #View the obs per set

y_test2.ravel()

```

```

[201]: array([1., 0., 0., 1., 1., 1., 1., 0., 1., 1., 1., 1., 1., 0., 1., 1., 0.,
        1., 1., 1., 1., 0., 1., 0., 1., 0., 1., 1., 0., 0., 1.,
        1., 1., 1., 1., 0., 1., 0., 1., 1., 0., 1., 0., 0., 1., 1., 1.,
        1., 0., 1., 0., 0., 1., 1., 1., 1., 1., 0., 1., 0., 1., 0., 1., 1.,
        1., 0., 1., 0., 1., 1., 1., 1., 0., 0., 0., 0.])

```

5 Build the model the models

5.1 Random Forest

```
[202]: from sklearn.ensemble import RandomForestClassifier
```

```
[203]: # Use the Random Forest Classifier Model

# Create random forest classifier object
randomforest = RandomForestClassifier(random_state=0, n_jobs=-1,
    ↪class_weight="balanced")

#Train model
rf_model = randomforest.fit(X_train2, y_train2.ravel())
```

```
[204]: # view model accuracy score
rf_model.score(X_test2, y_test2.ravel())
```

```
[204]: 0.9625
```

```
[205]: # Random Forest appears to have a reasonable accuracy
#
```

```
[206]: # Compute predictions

from sklearn.metrics import RocCurveDisplay, roc_auc_score,
    ↪average_precision_score, PrecisionRecallDisplay, ConfusionMatrixDisplay, \
    precision_recall_curve, accuracy_score
```

```
[207]: # compute predict probabilities on test set
# only need the probabilities for the positive class only

y_pred_probs = rf_model.predict_proba(X_test2)[:, 1]

# compute predictions on test set
y_preds = rf_model.predict(X_test2)
```

```
[208]: #Check the accuracy
print(f'Accuracy:, {accuracy_score(y_test2.ravel(), y_preds)}')
```

```
Accuracy:, 0.9625
```

```
[209]: # Compute AUC to validate accuracy
auc_score = roc_auc_score(y_test2.ravel(), y_pred_probs)
round(auc_score,4)
```

```
[209]: 0.999
```




```
[ ]: rf_model_new.score(X_test2_new, y_test2_new.ravel())
```

```
[210]: # Plot the ROC Curve
```

```
def roc_curve_plt(test, preds, Name, Title):  
    lw = 1  
    RocCurveDisplay.from_predictions(test, preds,  
                                     name = Name,  
                                     color="red", pos_label=None)  
    plt.plot([0, 1], [0, 1], 'k--', lw=lw)  
    plt.title(Title)  
    plt.show()  
    return
```

```
[211]: #Define parameters for ROC Curve plot
```

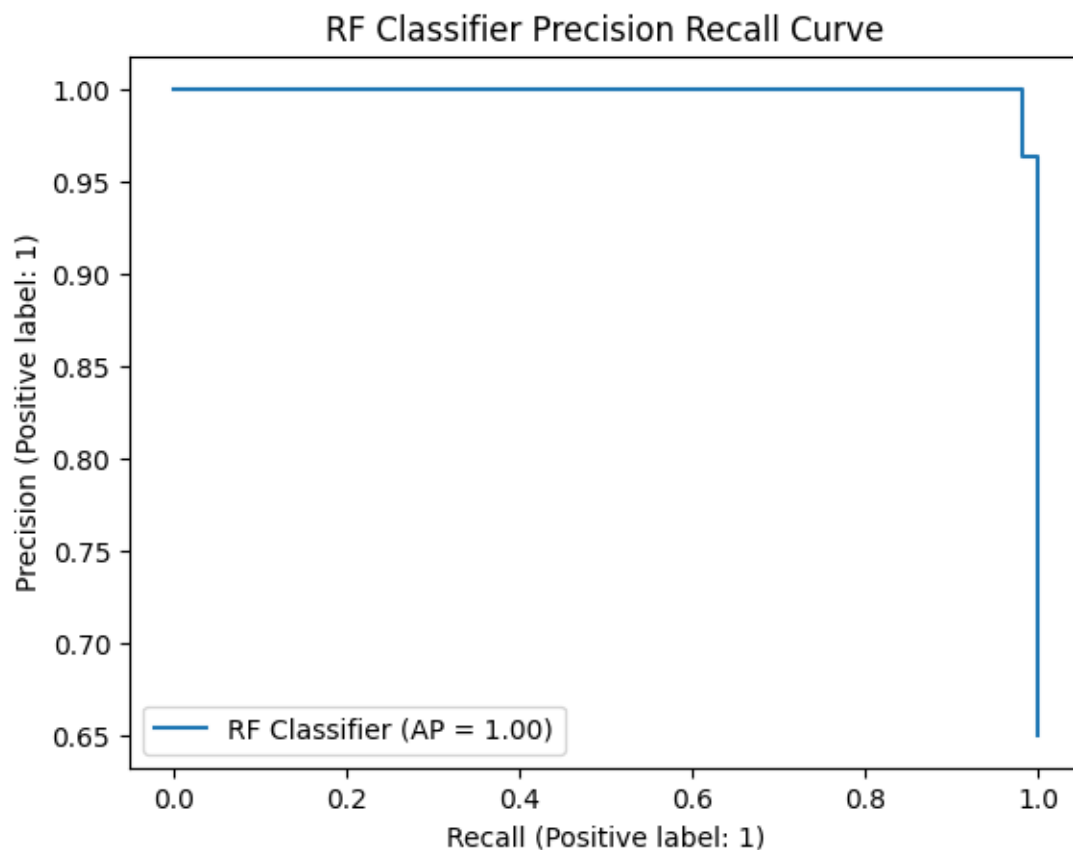
```
test = y_test2  
preds = y_preds  
Name = 'Random Forest Classifier'  
Title = 'Random Forest Classifier ROC Curve'  
roc_curve_plt(test, preds, Name, Title)
```



```
[212]: # Compute average precision score
# Plot the precision recall curve

def prc_plt(test, pred_probs, Name, Title):
    PrecisionRecallDisplay.from_predictions(
        test, pred_probs, name = Name)
    plt.title(Title)
    plt.show()
    return
```

```
[213]: #Define parameters for prc
test = y_test2
pred_probs = y_pred_probs
Name = "RF Classifier"
Title = "RF Classifier Precision Recall Curve"
prc_plt(test, pred_probs, Name, Title)
```



```
[216]: # Create Confusion Matrix from predictions
# Plot predictions on the test set
from sklearn.metrics import confusion_matrix, classification_report
```

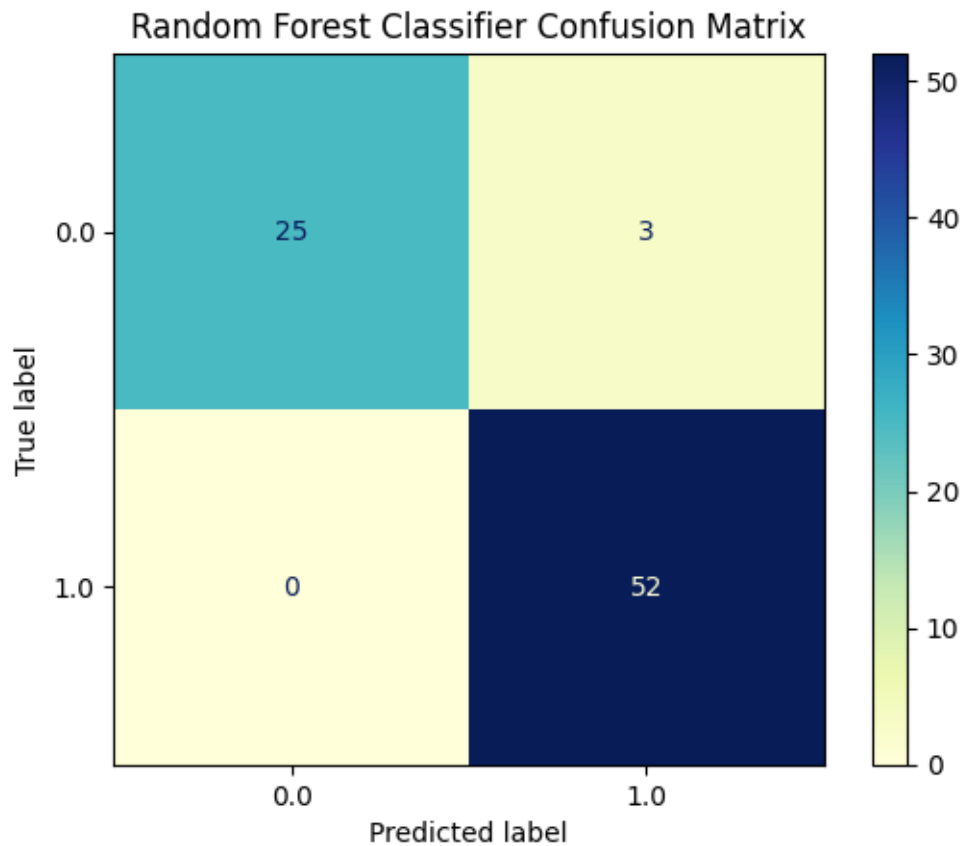
```

from sklearn.metrics import ConfusionMatrixDisplay

ConfusionMatrixDisplay.from_predictions(y_test2, y_preds, cmap= 'YlGnBu')

plt.title("Random Forest Classifier Confusion Matrix")
plt.show()

```



```

[217]: # Build Classification report
print(classification_report(y_test2, y_preds))

```

	precision	recall	f1-score	support
0.0	1.00	0.89	0.94	28
1.0	0.95	1.00	0.97	52
accuracy			0.96	80
macro avg	0.97	0.95	0.96	80
weighted avg	0.96	0.96	0.96	80

5.2 Build Adaboost model

[]:

```
[218]: from sklearn.ensemble import AdaBoostClassifier
       from sklearn.model_selection import cross_val_score
```

```
[219]: # Create adaboost tree classifier object
       adaboost = AdaBoostClassifier(random_state=0)

       # Train model on the important features data set
       ab_model = adaboost.fit(X_train2, y_train2.ravel())
```

```
[220]: # Cross-validate Adaboost model

       cvScore = cross_val_score(ab_model, X_test2, y_test2.ravel(),
                                   ↪scoring="accuracy")
       print("CV Scores:", cvScore)
       print("CV Scores Mean:", cvScore.mean())
```

```
CV Scores: [0.875  0.875  1.      0.9375 1.      ]
CV Scores Mean: 0.9375
```

```
[224]: # Tune hyperparameters with Adaboost
       from sklearn.model_selection import GridSearchCV

       # Define hyperparameter candidates
       param_grid = {'n_estimators': [100, 200, 400, 600, 800, 1000],
                     'learning_rate': [0.001, 0.01, 0.1, 0.2, 0.5]}

       # Instantiate GridSearchCV model
       gs_ab = GridSearchCV(AdaBoostClassifier(), param_grid = param_grid)

       # Train model
       gs_ab_model = gs_ab.fit(X_train2, y_train2.ravel())

       # Cross-validate model
       cvScore2 = cross_val_score(gs_ab_model, X_test2, y_test2.ravel(),
                                   ↪scoring="accuracy")
       print("CV Scores:", cvScore2)
       print("CV Scores Mean:", cvScore2.mean())
```

```
CV Scores: [0.9375 0.875  0.9375 0.9375 0.9375]
CV Scores Mean: 0.925
```

```
[222]: # View best hyperparameters
       gs_ab_model.best_params_
```

```
[222]: {'learning_rate': 0.5, 'n_estimators': 100}
```

```
[225]: # Compute predictions

# compute predict probabilities on test set
# only need the probabilities for the positive class only
y_pred_probs_ab = gs_ab_model.predict_proba(X_test2)[:, 1]

# compute predictions on test set
y_preds_ab = gs_ab_model.predict(X_test2)
```

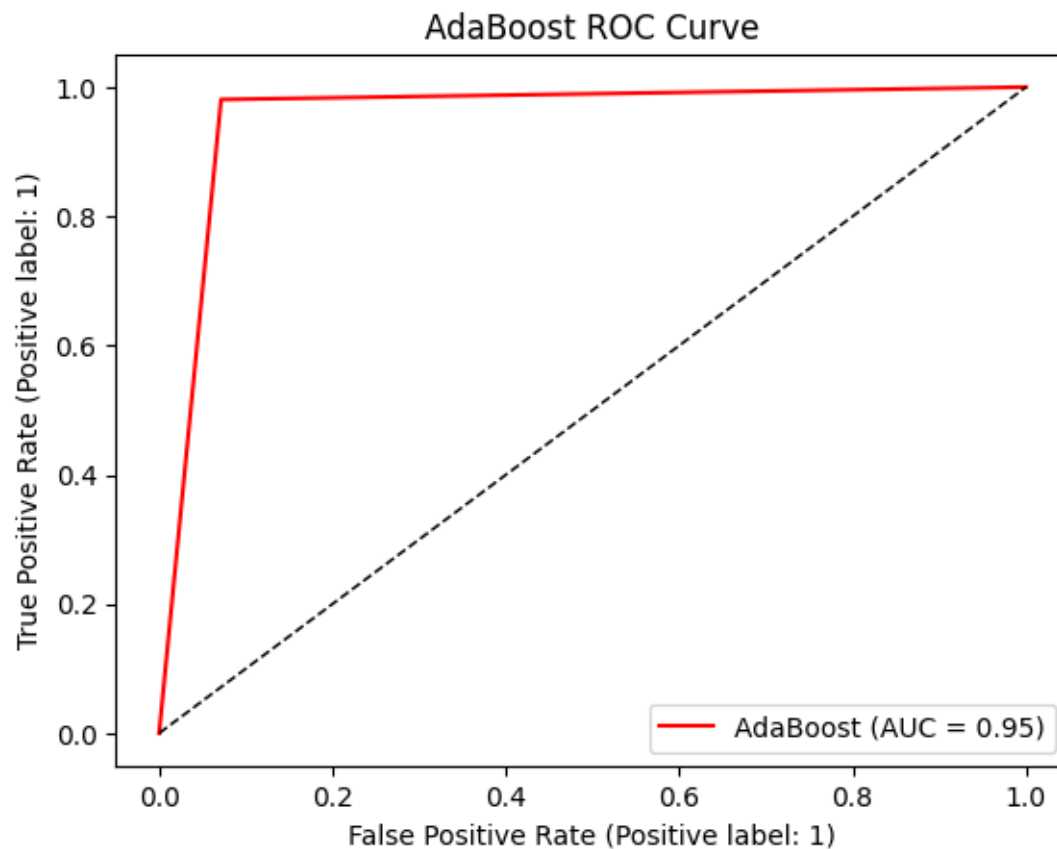
```
[226]: # Compute AUC
auc_score = roc_auc_score(y_test2, y_pred_probs_ab)
round(auc_score,4)
```

```
[226]: 0.9979
```

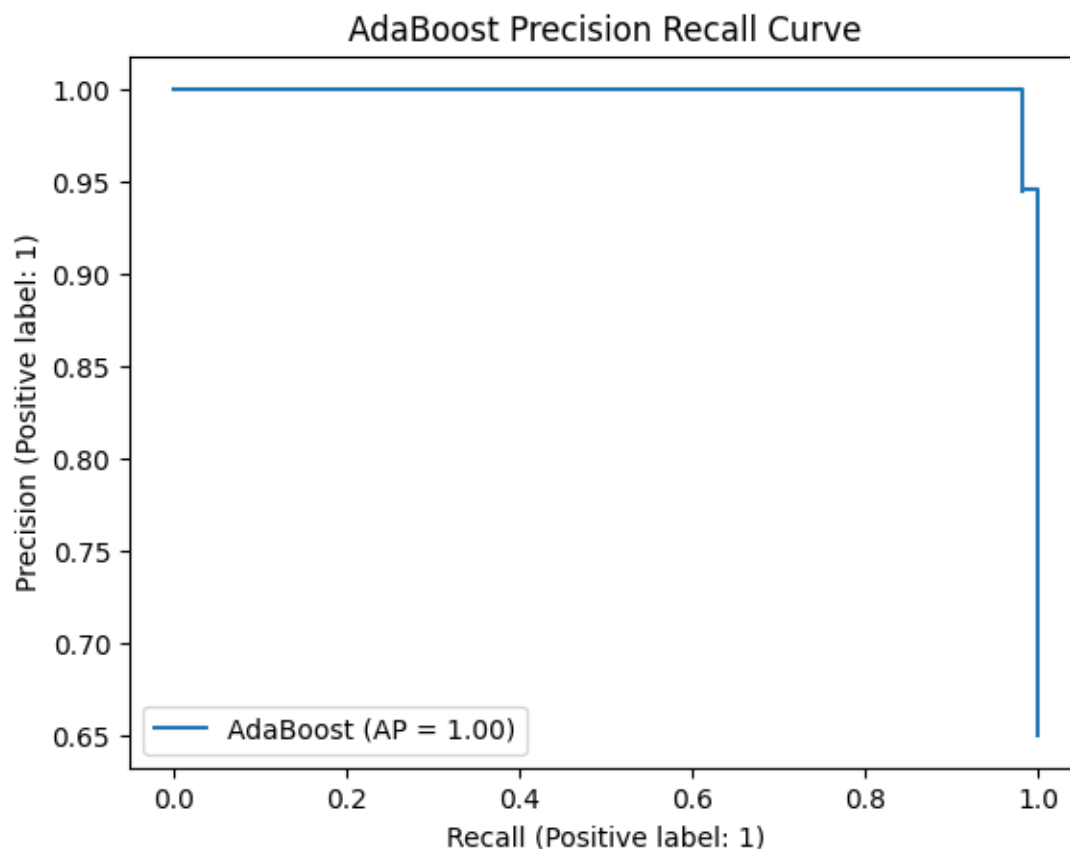
```
[227]: #Check the accuracy
print(f'Accuracy:, {accuracy_score(y_test2, y_preds_ab)}')
```

```
Accuracy:, 0.9625
```

```
[228]: #Define parameters for ROC Curve plot
test = y_test2
preds = y_preds_ab
Name = 'AdaBoost'
Title = 'AdaBoost ROC Curve'
roc_curve_plt(test, preds, Name,Title)
```



```
[229]: #Define parameters for prc  
test = y_test2  
pred_probs = y_pred_probs_ab  
Name = "AdaBoost"  
Title = "AdaBoost Precision Recall Curve"  
prc_plt(test, pred_probs, Name, Title)
```




```
[230]: # Build Classification report
print(classification_report(y_test2, y_preds_ab))
```

	precision	recall	f1-score	support
0.0	0.96	0.93	0.95	28
1.0	0.96	0.98	0.97	52
accuracy			0.96	80
macro avg	0.96	0.95	0.96	80
weighted avg	0.96	0.96	0.96	80

```
[231]: # Create Confusion Matrix from predictions

ConfusionMatrixDisplay.from_predictions(y_test2, y_preds_ab, cmap= 'YlGnBu')

plt.title("AdaBoost Confusion Matrix")
plt.show()
```



5.3 XgBoost Model

[250]: *# xgboost model*

```
import xgboost as xgb
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV
```

[]:

[251]: *#Define model parameters for input to the GridSearch*

```
params={'eta': [0.05, 0.10, 0.15, 0.25],
        'max_depth': [3, 4, 5, 6],
        'min_child_weight': [1, 3, 5],
        'gamma': [0.0, 0.1, 0.2, 0.4],
        'colsample_bytree': [0.3, 0.4, 0.7]}
```


[252]: *#Build and fit the model with cross-validation.*

```
xg = xgb.XGBClassifier()

grid = GridSearchCV(xg, params, n_jobs=4,
                    scoring="accuracy",
                    cv=3)

grid.fit(X_train2, y_train2)
```

[252]: GridSearchCV(cv=3,
 estimator=XGBClassifier(base_score=None, booster=None,
 callbacks=None, colsample_bylevel=None,
 colsample_bynode=None,
 colsample_bytree=None,
 early_stopping_rounds=None,
 enable_categorical=False, eval_metric=None,
 feature_types=None, gamma=None,
 gpu_id=None, grow_policy=None,
 importance_type=None,
 interaction_constraints=None,
 learning_rate=None,...
 max_delta_step=None, max_depth=None,
 max_leaves=None, min_child_weight=None,
 missing=nan, monotone_constraints=None,
 n_estimators=100, n_jobs=None,
 num_parallel_tree=None, predictor=None,
 random_state=None, ...),
 n_jobs=4,
 param_grid={'colsample_bytree': [0.3, 0.4, 0.7],
 'eta': [0.05, 0.1, 0.15, 0.25],
 'gamma': [0.0, 0.1, 0.2, 0.4],
 'max_depth': [3, 4, 5, 6],
 'min_child_weight': [1, 3, 5]},
 scoring='accuracy')

[253]: grid.best_estimator_

[253]: XGBClassifier(base_score=None, booster=None, callbacks=None,
 colsample_bylevel=None, colsample_bynode=None,
 colsample_bytree=0.3, early_stopping_rounds=None,
 enable_categorical=False, eta=0.25, eval_metric=None,
 feature_types=None, gamma=0.4, gpu_id=None, grow_policy=None,
 importance_type=None, interaction_constraints=None,
 learning_rate=None, max_bin=None, max_cat_threshold=None,
 max_cat_to_onehot=None, max_delta_step=None, max_depth=6,
 max_leaves=None, min_child_weight=1, missing=nan,

```
monotone_constraints=None, n_estimators=100, n_jobs=None,
num_parallel_tree=None, predictor=None, ...)
```

```
[254]: grid.best_params_
```

```
[254]: {'colsample_bytree': 0.3,
        'eta': 0.25,
        'gamma': 0.4,
        'max_depth': 6,
        'min_child_weight': 1}
```

```
[255]: best_params = {'eta':[0.25],
                      'max_depth':[6],
                      'min_child_weight': [1,3,5],
                      'gamma':[0.4],
                      'colsample_bytree':[0.3,]}
```

```
[256]: grid.n_features_in_
```

```
[256]: 13
```

```
[257]: #rerun the model with the best parameters
xg = xgb.XGBClassifier()

grid = GridSearchCV(xg, best_params, n_jobs=4,
                    scoring="accuracy",
                    cv=3)

grid.fit(X_train2, y_train2)
```

```
[257]: GridSearchCV(cv=3,
                  estimator=XGBClassifier(base_score=None, booster=None,
                                           callbacks=None, colsample_bylevel=None,
                                           colsample_bynode=None,
                                           colsample_bytree=None,
                                           early_stopping_rounds=None,
                                           enable_categorical=False, eval_metric=None,
                                           feature_types=None, gamma=None,
                                           gpu_id=None, grow_policy=None,
                                           importance_type=None,
                                           interaction_constraints=None,
                                           learning_rate=None, ...
                                           max_cat_to_onehot=None,
                                           max_delta_step=None, max_depth=None,
                                           max_leaves=None, min_child_weight=None,
                                           missing=nan, monotone_constraints=None,
                                           n_estimators=100, n_jobs=None,
```

```

num_parallel_tree=None, predictor=None,
random_state=None, ...),
n_jobs=4,
param_grid={'colsample_bytree': [0.3], 'eta': [0.25],
            'gamma': [0.4], 'max_depth': [6],
            'min_child_weight': [1, 3, 5]},
scoring='accuracy')

```

```

[258]: # Score the training set
grid.best_score_

```

```

[258]: 0.9749309351672251

```

```

[259]: #Score the test set
grid.score(X_test2, y_test2)

```

```

[259]: 0.975

```

```

[260]: #predictions
xg1_preds = grid.predict(X_test2)

```

```

[261]: # Compute predictions

# compute predict probabilities on test set
# only need the probabilities for the positive class only
pred_probs_xg1 = grid.predict_proba(X_test2)[:, 1]

```

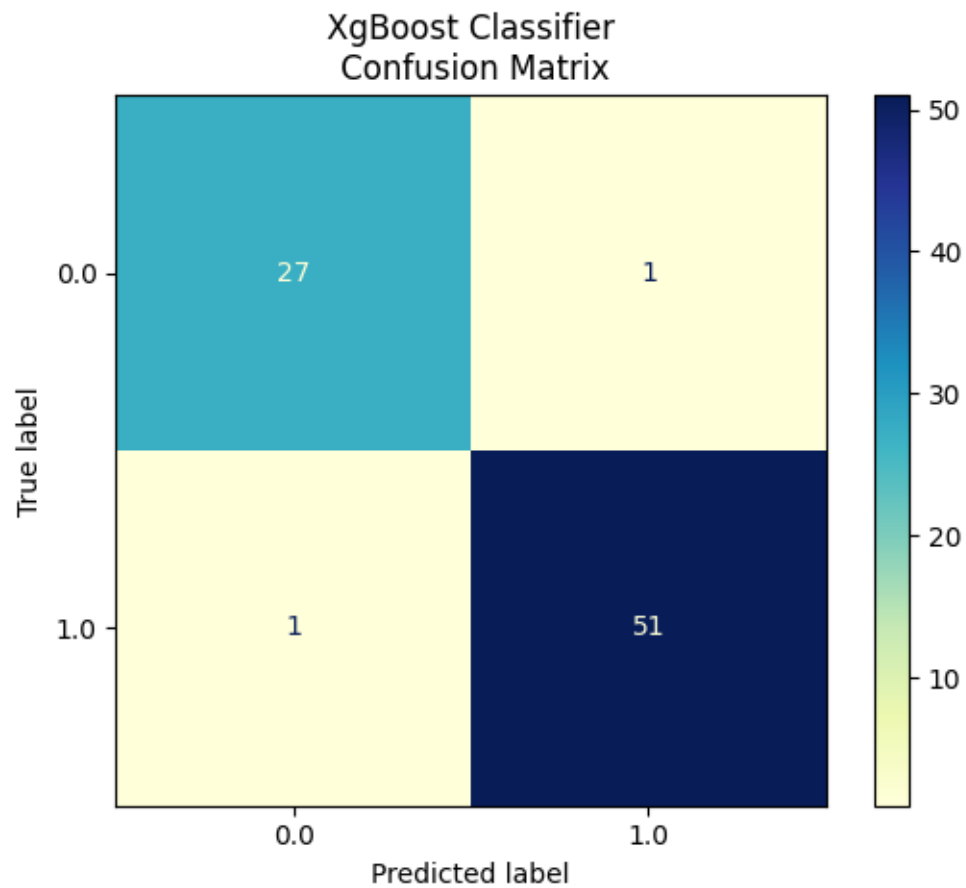
```

[262]: # Create Confusion Matrix from predictions

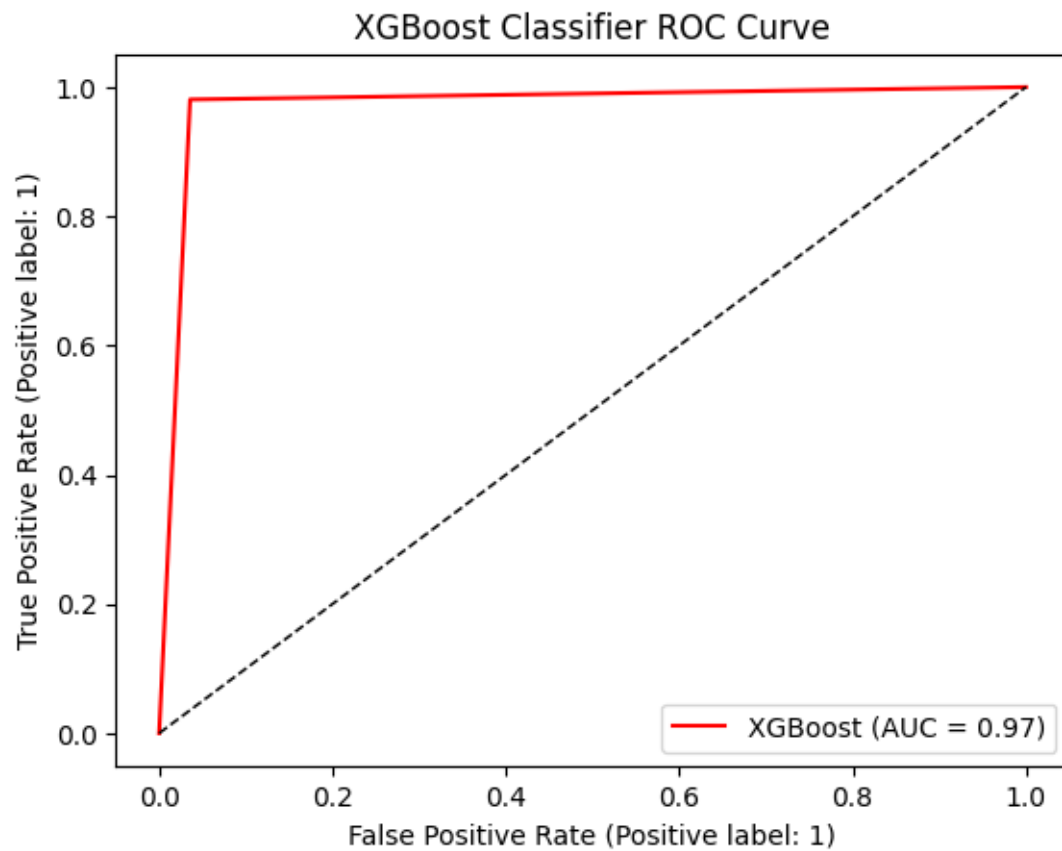
ConfusionMatrixDisplay.from_predictions(y_test2, xg1_preds, cmap= 'YlGnBu')

plt.title("XgBoost Classifier\n Confusion Matrix")
plt.show()


```



```
[263]: #Define parameters for ROC Curve plot  
test = y_test2  
preds = xg1_preds  
Name = 'XGBoost'  
Title = 'XGBoost Classifier ROC Curve'  
roc_curve_plt(test, preds, Name, Title)
```



```
[264]: #Define parameters for prc
test = y_test2
pred_probs = pred_probs_xg1
Name = "XGBoost"
Title = "XGBoost Precision Recall Curve"
prc_plt(test, pred_probs, Name, Title)
```



```
[265]: # Build Classification report
print(classification_report(y_test2, y_preds_ab))
```

	precision	recall	f1-score	support
0.0	0.96	0.93	0.95	28
1.0	0.96	0.98	0.97	52
accuracy			0.96	80
macro avg	0.96	0.95	0.96	80
weighted avg	0.96	0.96	0.96	80

6 Check the models against Validation Set

```
[236]: #Check the models against the validation set
#Compute predictions

def preds_counts(model):
```

```

val_preds = model.predict(val_scale_test) #val set obs

#create prediction dataframe
val_preds_df = pd.DataFrame(val_preds, columns=['class'])

#actual class
val_scale_class_df = pd.DataFrame(val_scale_class, columns=['actual_class'])

#add the preds column to the actuals df
val_scale_class_df['pred_class'] = val_preds_df['class']

# Compare the columns and display count
val_scale_class_df['diff'] = (val_scale_class_df.apply(lambda x:
↪x['actual_class'] if x['actual_class'] ==\
                                                    x['pred_class'] else
↪'False', axis =1))
# Count the differences
counts = val_scale_class_df.loc[val_scale_class_df['diff'] == 'False'].
↪value_counts()
return(counts, val_preds)

```

```

[233]: def roc_curve_plt(test, preds, Name, Title):
        lw = 1
        RocCurveDisplay.from_predictions(test, preds,
                                         name = Name,
                                         color="red", pos_label=None)

        plt.plot([0, 1], [0, 1], 'k--', lw=lw)
        plt.title(Title)
        plt.show()
        return

```

```

[234]: # Compute average precision score
# Plot the precision recall curve

def prc_plt(test, pred_probs, Name, Title):
    PrecisionRecallDisplay.from_predictions(
        test, pred_probs, name = Name)
    plt.title(Title)
    plt.show()
    return

```

6.1 RF with Validation Set

```

[266]: model = rf_model
        counts, rf_preds2 = preds_counts(model)
        counts

```

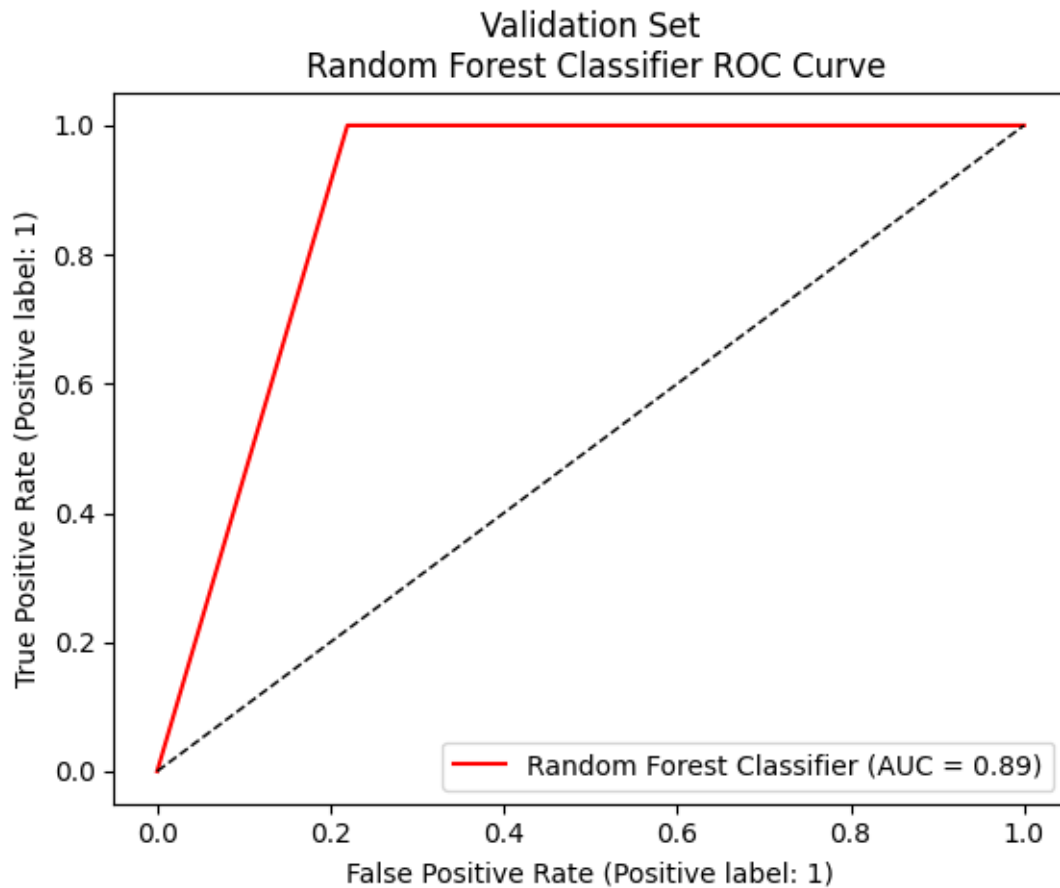
```
[266]: actual_class  pred_class  diff
      0.0          1.0        False    9
      dtype: int64
```

```
[268]: # Compute predictions

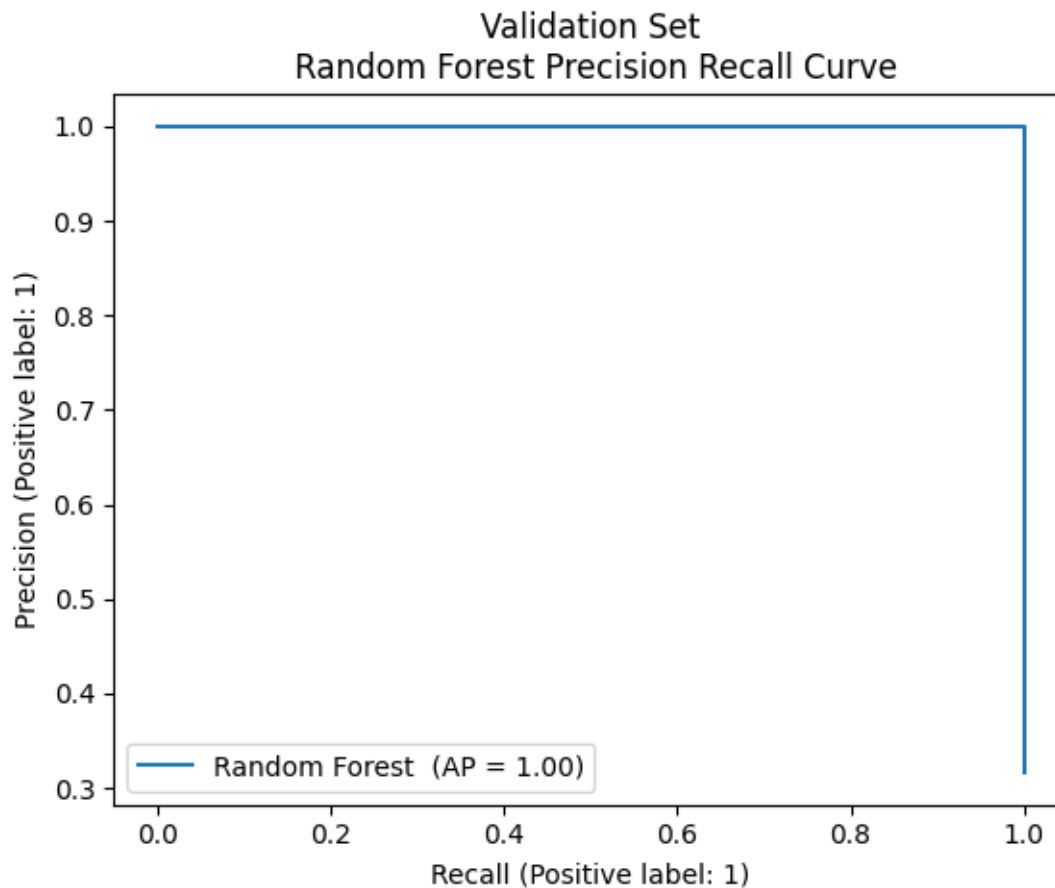
# compute predict probabilities on test set
# only need the probabilities for the positive class only
pred_probs_rf2 = rf_model.predict_proba(val_scale_test)[: , 1]
```

```
[ ]:
```

```
[269]: #Define parameters for ROC Curve plot
test = val_scale_class
preds = rf_preds2
Name = 'Random Forest Classifier'
Title = 'Validation Set\n Random Forest Classifier ROC Curve'
roc_curve_plt(test, preds, Name,Title)
```




```
[270]: #Define parameters for prc
test = val_scale_class
pred_probs = pred_probs_rf2
Name = "Random Forest "
Title = "Validation Set\n Random Forest Precision Recall Curve"
prc_plt(test, pred_probs, Name, Title)
```



6.2 AdaBoost with Validation Set

```
[ ]:
```

```
[237]: model = gs_ab_model
counts, ab_preds2 = preds_counts(model)
counts
```

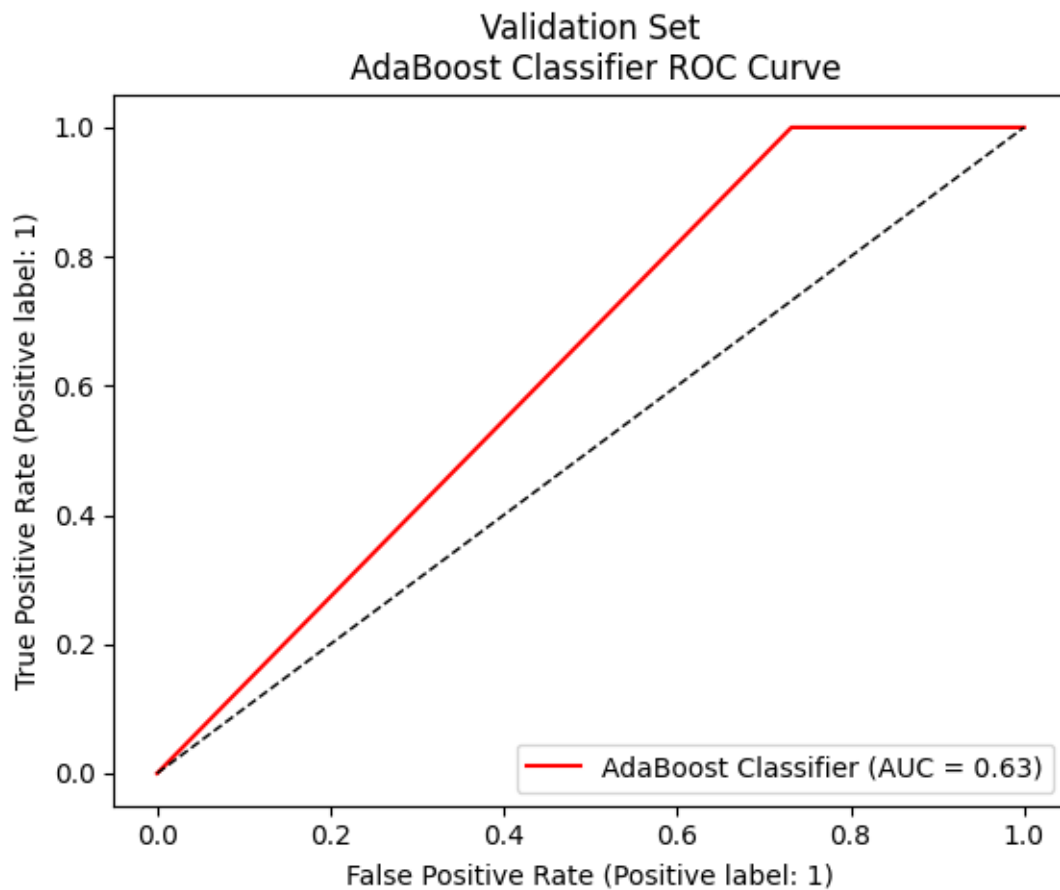
```
[237]: actual_class  pred_class  diff
0.0          1.0          False    30
dtype: int64
```

```
[239]: # Compute predictions

# compute predict probabilities on test set
# only need the probabilities for the positive class only
pred_probs_ab2 = gs_ab_model.predict_proba(val_scale_test)[: , 1]
```

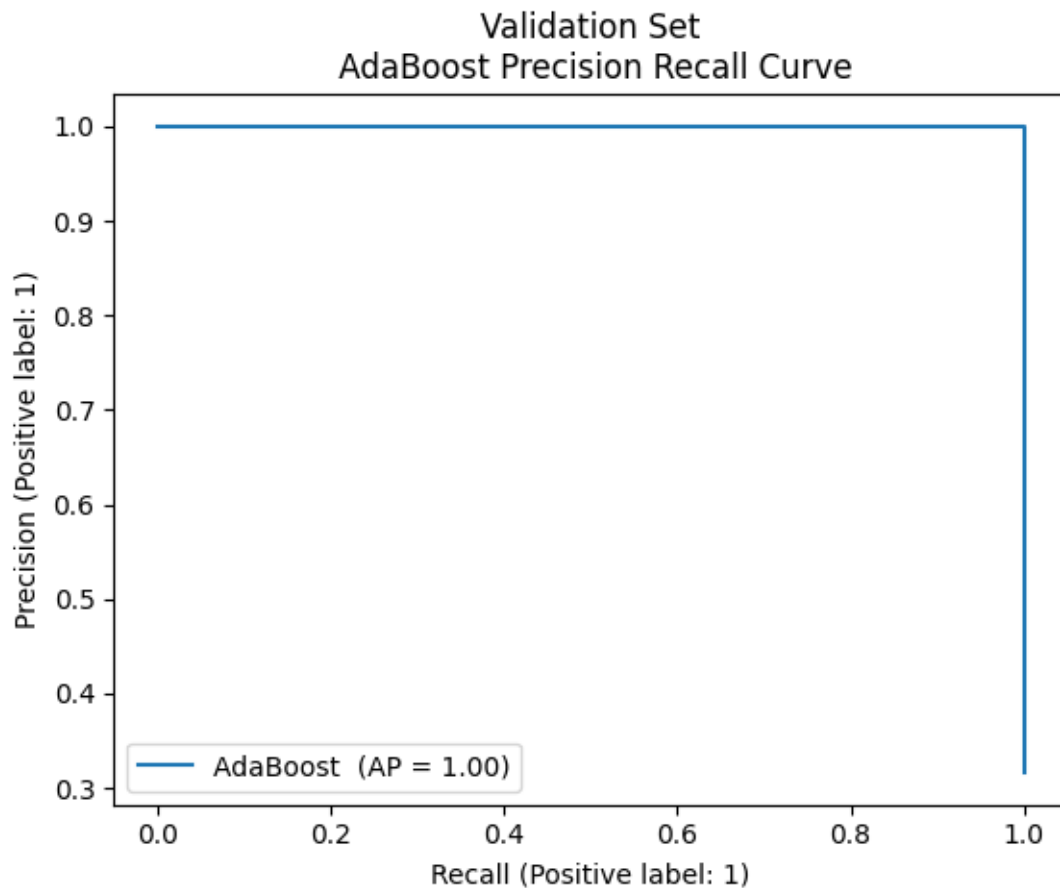
```
[ ]:
```

```
[242]: #Define parameters for ROC Curve plot
test = val_scale_class
preds = ab_preds2
Name = 'AdaBoost Classifier'
Title = 'Validation Set\n AdaBoost Classifier ROC Curve'
roc_curve_plt(test, preds, Name,Title)
```



```
[243]: #Define parameters for prc
test = val_scale_class
pred_probs = pred_probs_ab2
Name = "AdaBoost "
```

```
Title = "Validation Set\n AdaBoost Precision Recall Curve"
prc_plt(test, pred_probs, Name, Title)
```



6.3 XgBoost with Validation Set

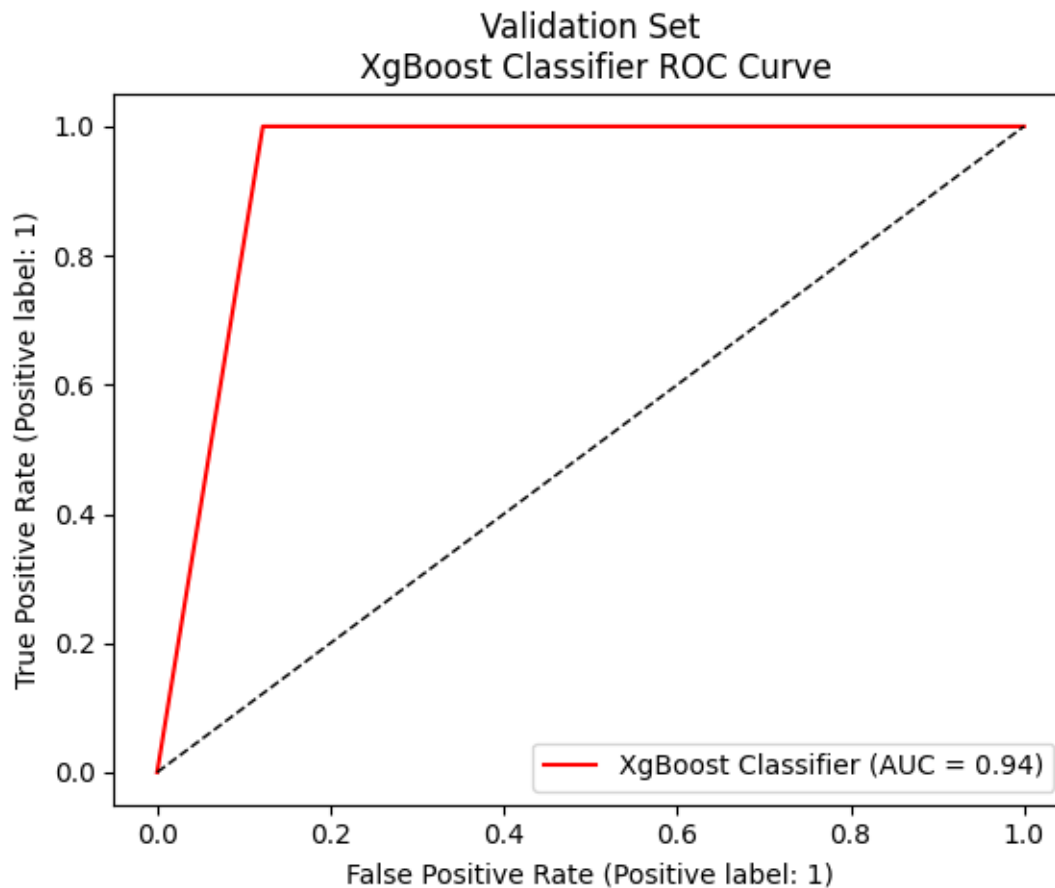
```
[271]: model = grid
counts, xg_preds2 = preds_counts(model)
counts
```

```
[271]: actual_class  pred_class  diff
0.0          1          False    5
dtype: int64
```

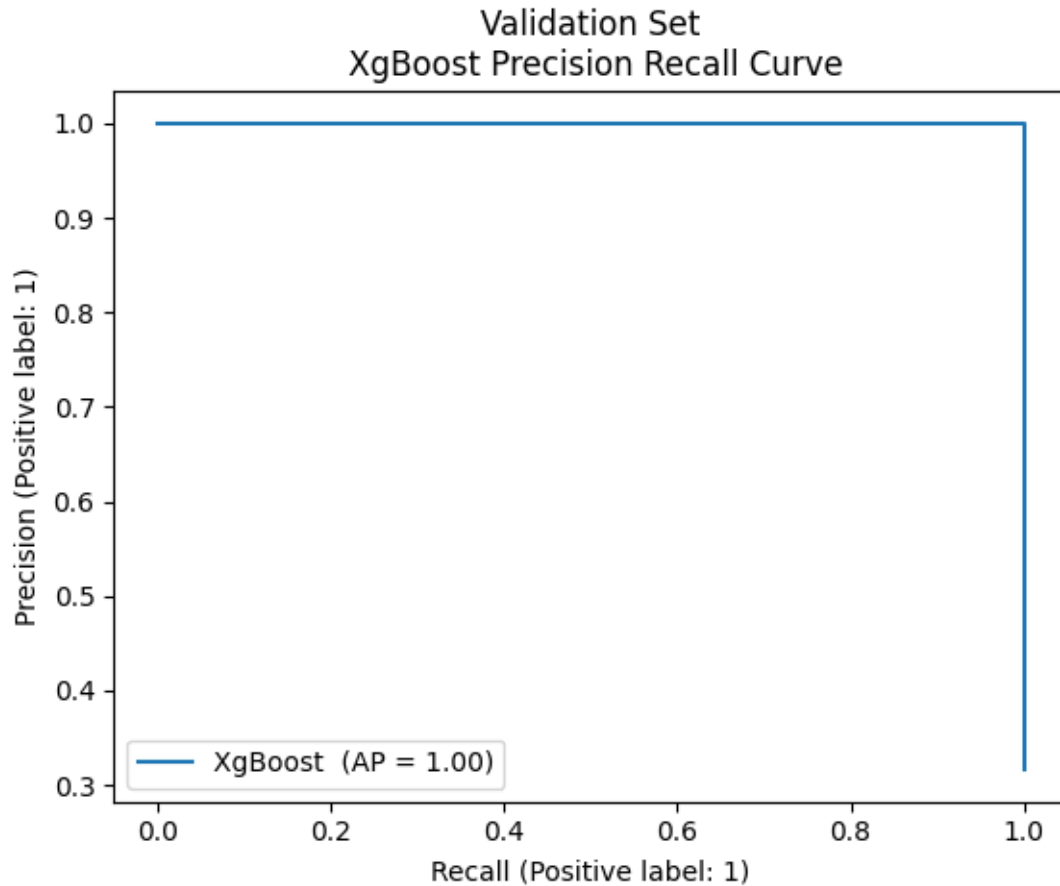
```
[272]: # Compute predictions

# compute predict probabilities on test set
# only need the probabilities for the positive class only
pred_probs_xg2 = grid.predict_proba(val_scale_test)[: , 1]
```

```
[274]: #Define parameters for ROC Curve plot
test = val_scale_class
preds = xg_preds2
Name = 'XgBoost Classifier'
Title = 'Validation Set\n XgBoost Classifier ROC Curve'
roc_curve_plt(test, preds, Name, Title)
```



```
[275]: #Define parameters for prc
test = val_scale_class
pred_probs = pred_probs_xg2
Name = "XgBoost "
Title = "Validation Set\n XgBoost Precision Recall Curve"
prc_plt(test, pred_probs, Name, Title)
```



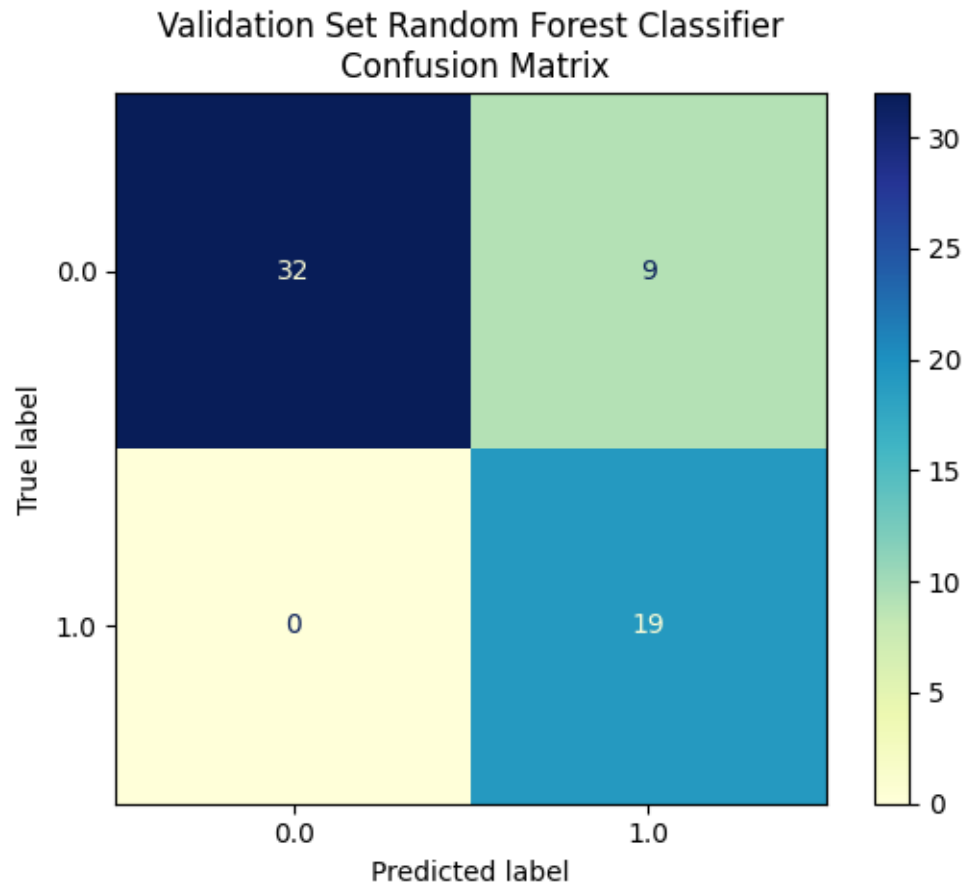
6.4 Confusion Matrices and Accuracy Scores

```
[276]: # Create Confusion Matrix from predictions using validation set
#val_scale = scaler.fit_transform(val_set_test)
#val_class_scale = scaler.fit_transform(val_set_class)

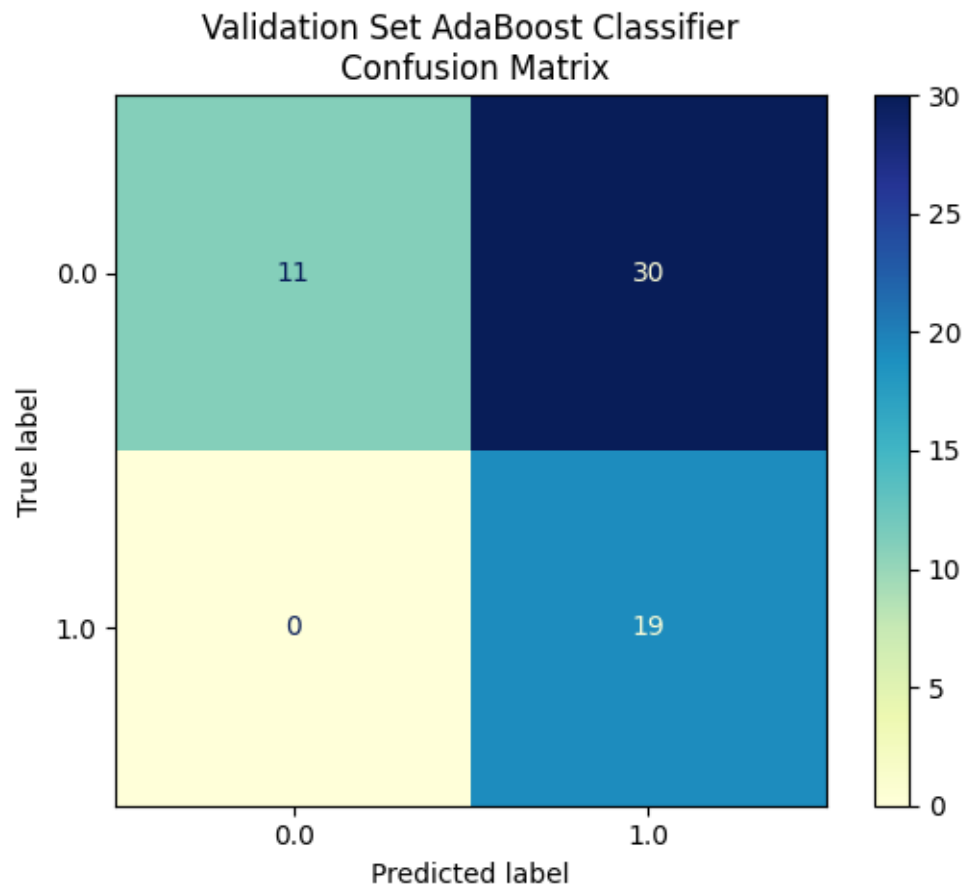
def cm_model(Title, preds):
    ConfusionMatrixDisplay.from_predictions(val_scale_class, preds, cmap=
↪ 'YlGnBu')

    plt.title(Title)
    plt.show()
    return
```

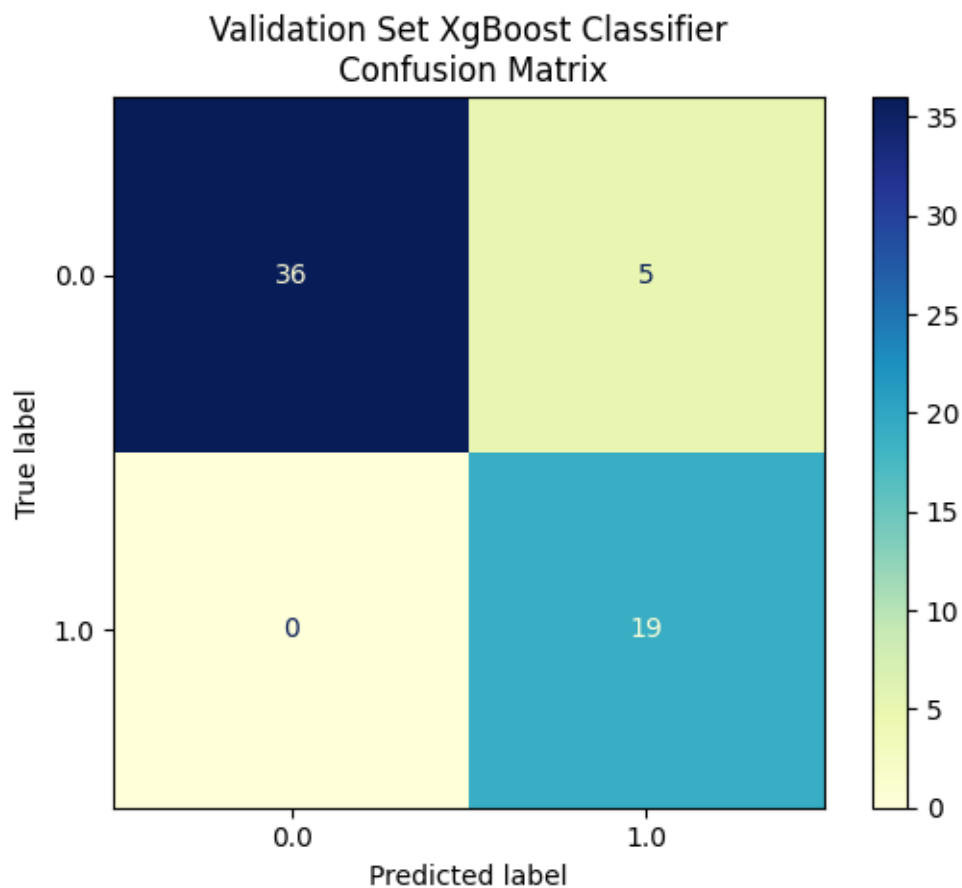
```
[277]: Title = 'Validation Set Random Forest Classifier\n Confusion Matrix'
preds = rf_preds2
cm_model(Title, preds)
```



```
[246]: # deinitely appears that AdaBoost suffers from overfitting  
  
Title = 'Validation Set AdaBoost Classifier\n Confusion Matrix'  
preds = ab_preds2  
cm_model(Title, preds)
```



```
[278]: Title = 'Validation Set XgBoost Classifier\n Confusion Matrix'  
preds = xg_preds2  
cm_model(Title, preds)
```



```
[248]: # Accuracy scores
# View accuracy score
def acc_score(preds):
    score = accuracy_score(val_scale_class, preds)
    return(score)
```

```
[279]: preds = rf_preds2
accuracy = acc_score(preds)
accuracy
```

[279]: 0.85

```
[249]: preds = ab_preds2
accuracy = acc_score(preds)
accuracy
```

[249]: 0.5


```
[280]: preds = xg_preds2  
accuracy = acc_score(preds)  
accuracy
```

```
[280]: 0.9166666666666666
```

```
[ ]:
```